

# DSCI 565: LINEAR NEURAL NETWORKS FOR CLASSIFICATION

Ke-Thia Yao

Lecture 5: 2025 September 10

2

## More on Generalization

# Error Measurements and the Test Set

3

- Given a dataset  $\mathcal{D}$ , empirical error  $\epsilon_{\mathcal{D}}$  computes the fraction of wrong classification for model  $f$ :

$$\epsilon_{\mathcal{D}}(f) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}(f(\mathbf{x}^{(i)}) \neq y^{(i)})$$

*indicator function*

- Population error  $\epsilon$  is the expected error fraction with some underlying population distribution  $P(X, Y)$  with probability density function  $p(x, y)$ :

$$\epsilon(f) = E_{(\mathbf{x}, y) \sim P} \mathbf{1}(f(\mathbf{x}) \neq y) = \int \int \mathbf{1}(f(\mathbf{x}) \neq y) p(\mathbf{x}, y) d\mathbf{x} dy$$

- We care about population error. Can approximate using empirical error, if
  - ▣  $\mathcal{D}$  is not used during training, i.e.,  $\mathcal{D}$  is the **test set**
  - ▣  $\mathcal{D}$  is sampled from the population distribution

# Central Limit Theorem

4

- Central limit theorem
  - ▣ Given any distribution with mean  $\mu$  and standard deviation  $\sigma$ , and  $n$  random samples from this distribution
  - ▣ As  $n$  approaches infinity the sample mean  $\hat{\mu}$  approaches  $\mu$  with sample standard deviation  $\sigma/\sqrt{n}$
- Treat each instance in the test set as a sample
- With large enough  $n$  the sample mean approaches the population mean
- If we want to reduce the sample standard deviation by  $1/2$ , then we must increase  $n$  by a factor of 4

# Test Set Size

5

- Probabilistically test error (0 or 1) of a random sample is just a Bernoulli random variable
- Variance of a Bernoulli random variable is  $p(1 - p)$  with maximum of 0.25, when  $p = 0.5$
- Setting aside the fact central limit theorem requires  $n$  to approach  $\infty$ 
  - ▣ If we want to be 68% confident (one std dev) sample mean is within  $\pm 0.01$  of population mean, then need 2500 samples  $\left(\sqrt{0.25/2500} = 0.01\right)$
  - ▣ If we want to be 95% confident (two std dev), then need 10,000 samples  $\left(\sqrt{0.25/10000} = 0.005\right)$

# Test Set Reuse

6

- In the strictest sense, a test set should only be used **once**
- Adaptive overfitting
  - ▣ Suppose you develop what you thought was the best model  $f_1$ , but it did not perform well on the test set
  - ▣ You go back develop  $f_2$
  - ▣ Then  $f_3$ , and so on
  - ▣ In a sense you have become part of machine learning algorithm, and you have seen the test set

# Statistical Learning Theory

7

- Instead of a test set, can we derive theoretical bounds on the generalization gap between empirical error and population error?
- Yes, *statistical learning theory* has developed such bounds
- These bound formulas are based on
  - ▣ The complexity (flexibility) of the model language, i.e., the VC dimension
  - ▣ Approximate correctness of the model  $\alpha$ , i.e., the generalization gap
  - ▣ The probability  $\delta$  of not finding an approximate correct model
  - ▣ The number of training examples  $n$
- For example, we may want  $\alpha = 0.1$  and  $\delta = 0.2$ , what should be the size training example  $n$ ?

# Bounding the Generalization Gap

8

Approximately Correct Model

$$P \left( \underbrace{R[p, f] - R_{\text{emp}}[\mathbf{X}, \mathbf{Y}, f]}_{\text{Generalization Gap}} < \alpha \right) \geq 1 - \delta \text{ for } \alpha \geq c \sqrt{(\text{VC} - \log \delta)/n}.$$

Probably Approximately Correct Model

- For a given modeling language ( $VC$ ) and lost function constant ( $c$ ), to satisfy the inequality, we can
  - ▣ Get more training example,  $\mathcal{O}(1/\sqrt{n})$
  - ▣ Allow less correct models, increase  $\alpha$
  - ▣ Lower probability of finding approximately correct model, increase  $\delta$



# Vapnik Chervonenkis Dimension


9

- Vapnik-Chervonenkis (VC) dimension measures the representational power of the modeling language  $\mathcal{H}$
- VC dimension of  $\mathcal{H}$  is the number of points  $\mathcal{H}$  can **shatter**
- Shattering  $n$  points
  - ▣ A set on  $n$  points can be labelled  $2^n$  ways using binary labels  $+, -$
  - ▣ If there exists a model  $h \in \mathcal{H}$  for every possible  $2^n$  labelling of **a set** of  $n$  points, then the VC dimension of  $\mathcal{H}$  is at least  $n$
- *Note: does not have to shatter all possible sets of  $n$  points, just have to shatter as least one set of  $n$  points*

# VC Dimension of Linear Models

10

- For 2-dimensional inputs, VC dimension is 3



[https://en.wikipedia.org/wiki/Vapnik%E2%80%93Chervonenkis\\_dimension](https://en.wikipedia.org/wiki/Vapnik%E2%80%93Chervonenkis_dimension)

- For  $d$ -dimensional inputs, VC dimension is  $d+1$

# VC Dimension of Neural Networks

11

- Given a neural network with sigmoid activation function of  $\|V\|$  nodes and  $\|E\|$  edges
  - ▣ VC dimension is at least  $\Omega(\|E\|^2)$
  - ▣ VC dimension is at most  $O(\|E\|^2\|V\|^2)$

# Deployment Environment

15

- ❑ Previous analysis assumes the training set is sampled from the population distribution
- ❑ But, if the actual environment in which we deploy the model may be different then the assumed population distribution
- ❑ Or perhaps the environment changes over time (patient population ages)
- ❑ Or the environment adapts to our model (spam detection)
- ❑ These are all causes of ***distribution shift***

# Types of Distribution Shift: Covariate Shift

16

- Assume our training data is sampled from  $P_S(\mathbf{x}, y)$ , but the test data is sampled from  $P_T(\mathbf{x}, y)$
- Covariate Shift assumes  $P(\mathbf{x})$  changes, but  $P(y|\mathbf{x})$  does not change

cat



cat



dog



dog



cat



cat



dog



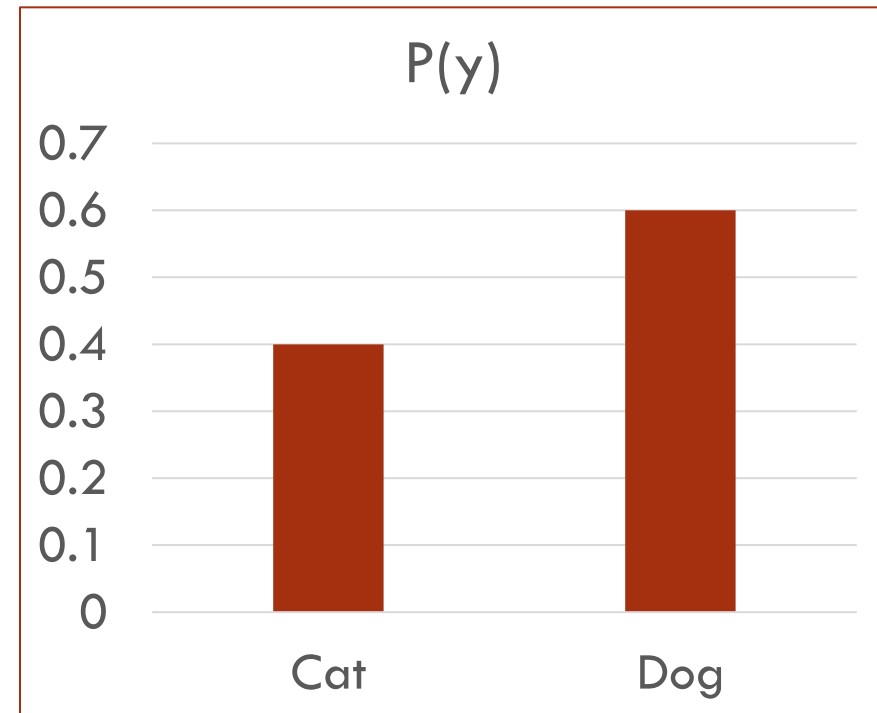
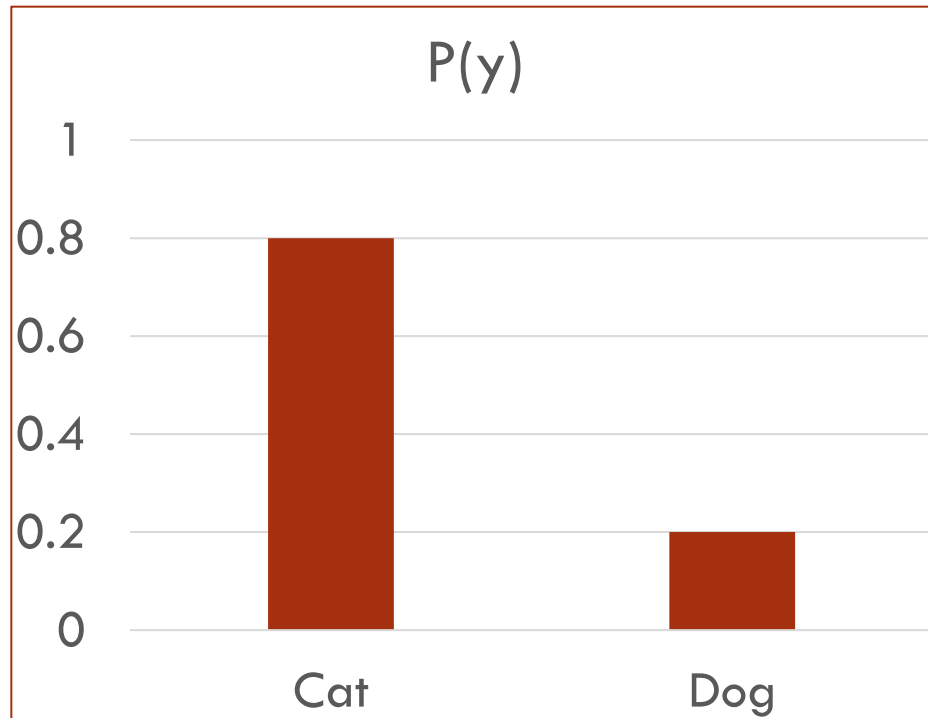
dog



# Types of Distribution Shift: Label Shift

17

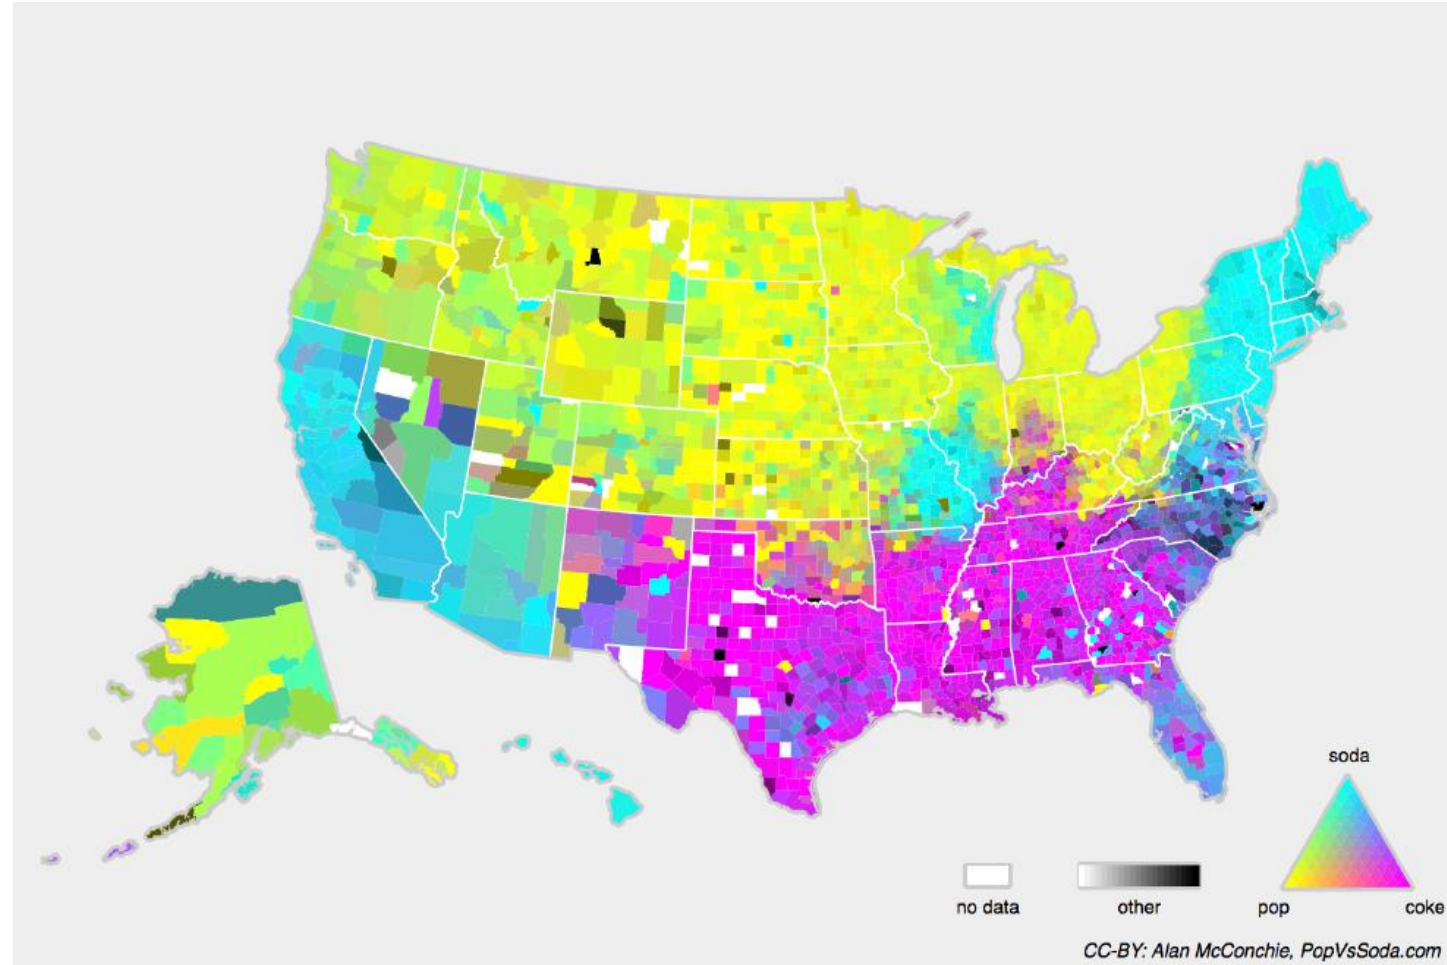
- Label shift assumes  $P(y)$  changes, but  $P(x|y)$  does not change
- Label shift is sometimes called prior probability shift



# Types of Distribution Shift: Concept Shift

18

- The definition of the label changes, i.e.,  $P(y|x)$
- For example, label for soft drink differs across US



# Examples of Distribution Shift

19

- Medical diagnostics: cancer detector works wonderfully on train/test set, but fails miserably on deployment
  - ▣ Oncologist did not provide enough negative instances. Additionally negatives instances gathered from university student volunteers.
- Object detection: tank detector failed to tanks in in forest
  - ▣ Images without tanks were taken in morning, and images with tanks were taken later in the day
- Recommendation system: continuous to recommend Santa hats long after Christmas
- ...



# Covariate Shift Correction

20

- The loss over the true population  $p(\mathbf{x}, y)$  is

$$E_{p(\mathbf{x}, y)} [l(f(\mathbf{x}), y)] = \int \int l(f(\mathbf{x}), y) p(\mathbf{x}, y) d\mathbf{x} dy.$$

- But training distribution is drawn from  $q(\mathbf{x}, y)$  with  $p(y|\mathbf{x}) = q(y|\mathbf{x})$
- We write population loss as

$$\int \int l(f(\mathbf{x}), y) p(y | \mathbf{x}) p(\mathbf{x}) d\mathbf{x} dy = \int \int l(f(\mathbf{x}), y) q(y | \mathbf{x}) q(\mathbf{x}) \boxed{\frac{p(\mathbf{x})}{q(\mathbf{x})}} d\mathbf{x} dy$$

Adjust weight of each instance

# Covariate Shift Correction

21

- If we know the ratio:

$$\beta_i \stackrel{\text{def}}{=} \frac{p(\mathbf{x}_i)}{q(\mathbf{x}_i)}$$

- Then we train model to minimize weighted empirical risk:

$$\underset{f}{\text{minimize}} \frac{1}{n} \sum_{i=1}^n \beta_i l(f(\mathbf{x}_i), y_i).$$

- In practice we do not know  $\beta$ . But if we can sample from  $p(\mathbf{x})$  then we can learn  $\beta$ .

# Learning $\beta$ Correction Using Logistic Regression

22

- Suppose points drawn from  $p$  is labelled  $z = 1$ , and from  $q$  is labelled  $z = -1$  (our training data)

- Then

$$P(z = 1 \mid \mathbf{x}) = \frac{p(\mathbf{x})}{p(\mathbf{x}) + q(\mathbf{x})} \text{ and hence } \frac{P(z = 1 \mid \mathbf{x})}{P(z = -1 \mid \mathbf{x})} = \frac{p(\mathbf{x})}{q(\mathbf{x})}$$

- Using logistic regression:

$$P(z = 1 \mid \mathbf{x}) = \frac{1}{1 + \exp(-h(\mathbf{x}))}$$

- Then:

$$\beta_i = \frac{1 / (1 + \exp(-h(\mathbf{x}_i)))}{\exp(-h(\mathbf{x}_i)) / (1 + \exp(-h(\mathbf{x}_i)))} = \exp(h(\mathbf{x}_i))$$

# A Taxonomy of Learning Problems

23

- Batch Learning
  - ▣ Learn using entire labelled training data.
  - ▣ No distributional shift. Learn once, deploy, and never have to change.
- Online Learning
  - ▣ Learning one sample at a time  $(x_i, y_i)$
  - ▣ Continuously adjust model according to success of predicting  $y_i$
- Control
  - ▣ Learn actions to control the environment (e.g., control car air conditioner)
  - ▣ Adjust model based on the response environment (hot parked car → max AC)
- Reinforcement Learning
  - ▣ Learn how environment behaves and learn policies on how to act
  - ▣ Environment responds in complex ways, e.g., adversarial for competitive games (chess, Go) or cooperative (allowing autonomous car to change lanes)

# Fairness, Accountability and Transparency in Machine Learning

24

- Machine learning models are being used to guide decision making in the real world
- Before deploying a model
  - ▣ Analyze the impact of using model
  - ▣ Make sure its decisions are appropriate for various subpopulations
  - ▣ Setup a governance structure for monitoring and management
- COMPAS: a deployed ML system for criminal risk assessment
  - ▣ Systematically gave black defendants higher risk,
  - ▣ Of all defendants receiving higher risk black defendants have lower percentage of recidivism
- Runaway feedback
  - ▣ After reading a few webpages on a conspiracy theory, systems recommend more webpages about conspiracy theory
  - ▣ Predictive policing system repeatedly sends patrol to the same neighborhoods