

# DSCI 565: CONVOLUTIONAL NEURAL NETWORKS

*This content is protected and may not  
be shared, uploaded, or distributed.*

Ke-Thia Yao

Lecture 8: 2025-09-22

# Convolutional Neural Networks


2

- Image data has a grid structure of pixels
- The neural networks that we have seen so far requires flattening of images in into vectors, which destroys the grid structure
- Can we some how preserve this grid structure?


# Initial Attempt

3

- Give an input image  $\mathbf{X}$ , to preserve the grid structure we require the hidden layer  $\mathbf{H}$  to have same shape as  $\mathbf{X}$
- Let  $[\mathbf{X}]_{i,j}$  and  $[\mathbf{H}]_{i,j}$  denote the pixel at  $(i,j)$
- $[\mathbf{H}]_{i,j}$  captures some feature related to  $(i,j)$



$\mathbf{X}$



$\mathbf{H}$

# Initial Attempt

4

- Then, if we use a fully connect architecture as before we need
  - ▣ A second-order tensor for the bias  $[\mathbf{U}]_{i,j}$
  - ▣ A fourth-order tensor for the weights  $[\mathbf{W}]_{i,j,k,l}$
- Then, to compute  $[\mathbf{H}]_{i,j}$

$$[\mathbf{H}]_{i,j} = [\mathbf{U}]_{i,j} + \sum_k \sum_l [\mathbf{W}]_{i,j,k,l} [\mathbf{X}]_{k,l}$$

$(10^3)^4 = 10^{12}$

- For a megapixel image (1 000x1 000), the weight matrix would have  $10^{12}$  (a trillion parameters)
- The largest Large Language Models (LLMs) is about 200 billion parameters
- What else is wrong with this initial attempt?

# Translational Invariance

5

- Detecting Waldo at one part of the image should be no different than any other part


$$[\mathbf{H}]_{i,j} = [\mathbf{U}]_{i,j} + \sum_k \sum_l [\mathbf{W}]_{i,j,k,l} [\mathbf{X}]_{k,l}$$



# Locality

6

- To capture the feature at location  $(i, j)$  we should not have to look at the entire image
- Only look at nearby pixels  $(i + a, j + b)$  and  $a, b \in [-\Delta, +\Delta]$

$$\begin{aligned} [\mathbf{H}]_{i,j} &= [\mathbf{U}]_{i,j} + \sum_k \sum_l [\mathbf{W}]_{i,j,k,l} [\mathbf{X}]_{k,l} \\ &= [\mathbf{U}]_{i,j} + \sum_a \sum_b [\mathbf{V}]_{i,j,a,b} [\mathbf{X}]_{i+a,j+b} \end{aligned}$$


Translational  
Invariance

$$[\mathbf{H}]_{i,j} = u + \sum_a \sum_b [\mathbf{V}]_{a,b} [\mathbf{X}]_{i+a,j+b}$$

Locality

$$[\mathbf{H}]_{i,j} = u + \sum_{a=-\Delta}^{\Delta} \sum_{b=-\Delta}^{\Delta} [\mathbf{V}]_{a,b} [\mathbf{X}]_{i+a,j+b}$$

# Channels

7

- So far, we have ignored the fact that some images have color channels
- With color
  - ▣ input becomes a third-order tensor  $[\mathbf{X}]_{i,j,c}$  and
  - ▣ weight because a third-order tensor as well  $[\mathbf{V}]_{a,b,c}$
- Also, instead of just one hidden feature map we want multiple feature maps  $[\mathbf{H}]_{i,j,d}$ , and weight becomes  $[\mathbf{V}]_{a,b,c,d}$

$$[\mathbf{H}]_{i,j} = u + \sum_{a=-\Delta}^{\Delta} \sum_{b=-\Delta}^{\Delta} [\mathbf{V}]_{a,b} [\mathbf{X}]_{i+a,j+b}$$

$$[\mathbf{H}]_{i,j,d} = \sum_{a=-\Delta}^{\Delta} \sum_{b=-\Delta}^{\Delta} \sum_c [\mathbf{V}]_{a,b,c,d} [\mathbf{X}]_{i+a,j+b,c}$$

# Convolution vs Cross-Correlation

8

- In machine learning, we call the previous formulas convolution
- But, in mathematics it is called **cross-correlation**
- In mathematics, convolution direction is “flipped”

$$(f * g)(\mathbf{x}) = \int f(\mathbf{z})g(\mathbf{x} - \mathbf{z})d\mathbf{z}.$$



# Cross-Correlation Operation

9

- Example of a cross-correlation operation using a 2x2 kernel

Input		Kernel		Output																		
<table border="1"><tr><td>0</td><td>1</td><td>2</td></tr><tr><td>3</td><td>4</td><td>5</td></tr><tr><td>6</td><td>7</td><td>8</td></tr></table>	0	1	2	3	4	5	6	7	8	*	<table border="1"><tr><td>0</td><td>1</td></tr><tr><td>2</td><td>3</td></tr></table>	0	1	2	3	=	<table border="1"><tr><td>19</td><td>25</td></tr><tr><td>37</td><td>43</td></tr></table>	19	25	37	43	$\begin{aligned} 0 \times 0 + 1 \times 1 + 3 \times 2 + 4 \times 3 &= 19, \\ 1 \times 0 + 2 \times 1 + 4 \times 2 + 5 \times 3 &= 25, \\ 3 \times 0 + 4 \times 1 + 6 \times 2 + 7 \times 3 &= 37, \\ 4 \times 0 + 5 \times 1 + 7 \times 2 + 8 \times 3 &= 43. \end{aligned}$
0	1	2																				
3	4	5																				
6	7	8																				
0	1																					
2	3																					
19	25																					
37	43																					


- The resulting output shape reduce by 1 in each dimension
- In general, if the input size is  $n_h \times n_w$  and kernel size is  $k_h \times k_w$  output shape is

$$(n_h - k_h + 1, n_w - k_w + 1)$$

# Inspiration for Mammal Visual Cortex

10


- Cats (and monkeys) have many neurons in the visual cortex of the brain have small local receptive fields
- They only activate when a limited region of the visual field is stimulated
- Moreover, some of these neurons respond to lines in certain directions



*Figure 14-1. Biological neurons in the visual cortex respond to specific patterns in small regions of the visual field called receptive fields; as the visual signal makes its way through consecutive brain modules, neurons respond to more complex patterns in larger receptive fields*

# Convolution Filter

Use **convolution filter** to mimic neurons with limited receptive fields



Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

# Notebook




12

□ `chapter_convolutional-neural-networks/conv-layer.ipynb`

# Padding

13


- For convolutions
  - ▣ Pixels on edges of images are used less often
  - ▣ Size of the image shrinks
- Shrinking becomes problematic if we have multiple convolution layers



# Padding

14

- Padding adds extra zero pixels around the border



- With padding  $p_h$  and  $p_w$  the output shape is
$$(n_h - k_h + p_h + 1, n_w - k_w + p_w + 1)$$
$$= (3 - 2 + 2 + 1, 3 - 2 + 2 + 1) = (4, 4)$$

# Padding and Kernel Size

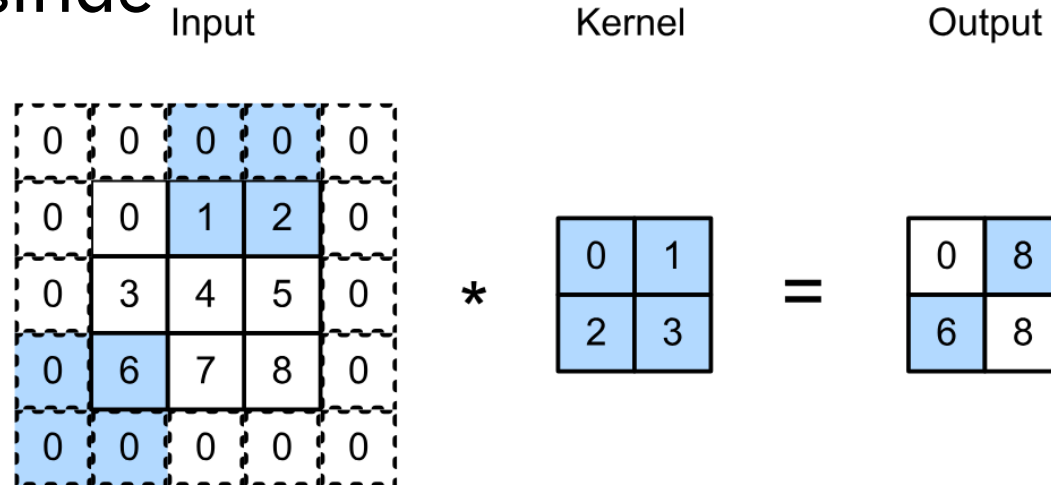
15

- To have the same number of pixel of each side of the border:  
 $p_h$  and  $p_w$  must be even
- To keep the same image size:  
 $p_h = k_h - 1$  and  $p_w = k_w - 1$
- Then kernel size must be odd
- Having the same image size makes the interpretation of the hidden layer simpler:  
 $[\mathbf{H}]_{i,j}$  is calculated using pixels centered around  $[\mathbf{X}]_{i,j}$
- Note: PyTorch padding `Conv2d(..., padding=p)` is  $p_h = p_w = 2p$

# Stride

16

- By default, the convolution slides one pixel at a time (stride of 1)
- To decrease computational complexity and/or to downsample, we can increase the stride



- With stride  $s_h$  and  $s_w$ , the output shape is

$$\lfloor (n_h - k_h + p_h + s_h) / s_h \rfloor \times \lfloor (n_w - k_w + p_w + s_w) / s_w \rfloor$$



# Notebook


17

□ `chapter_convolutional-neural-networks/padding-and-strides.ipynb`

# Multiple Input and Multiple Output Channels

18

- Three input channels (RGB colors) with five output channels



# of filters/channels/feature maps

- With input channel size  $c_i$ , the kernel shape becomes  $c_i \times k_h \times k_w$

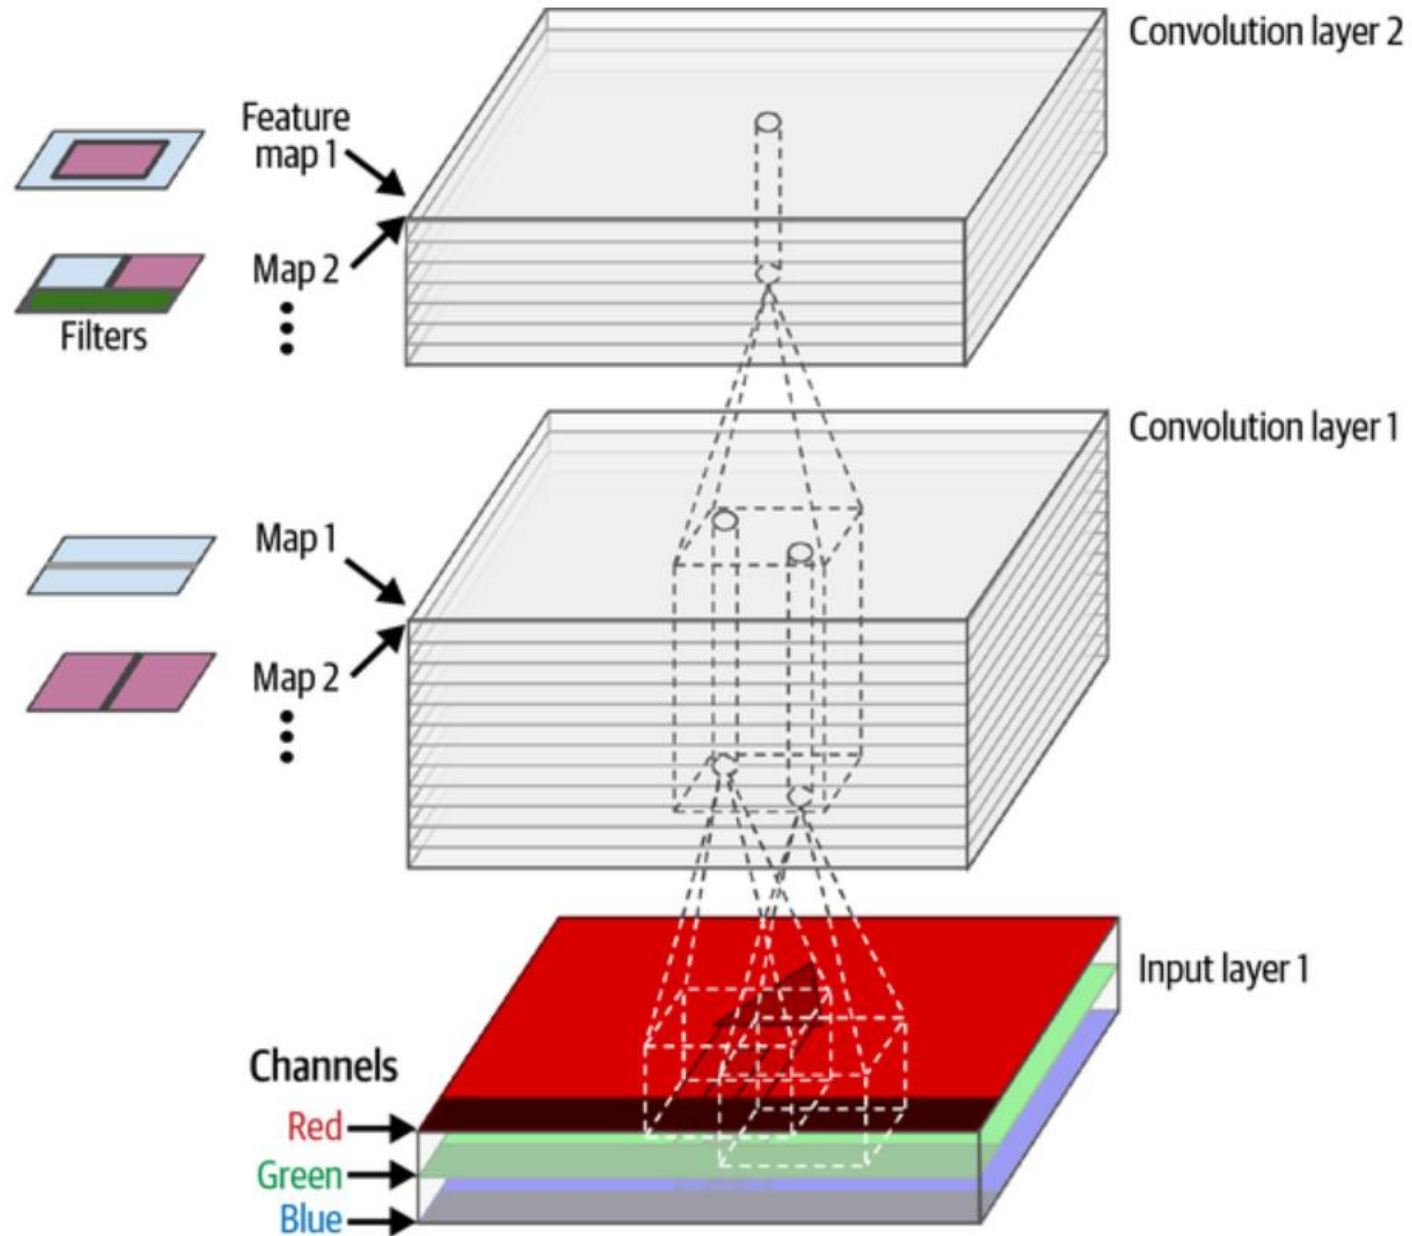
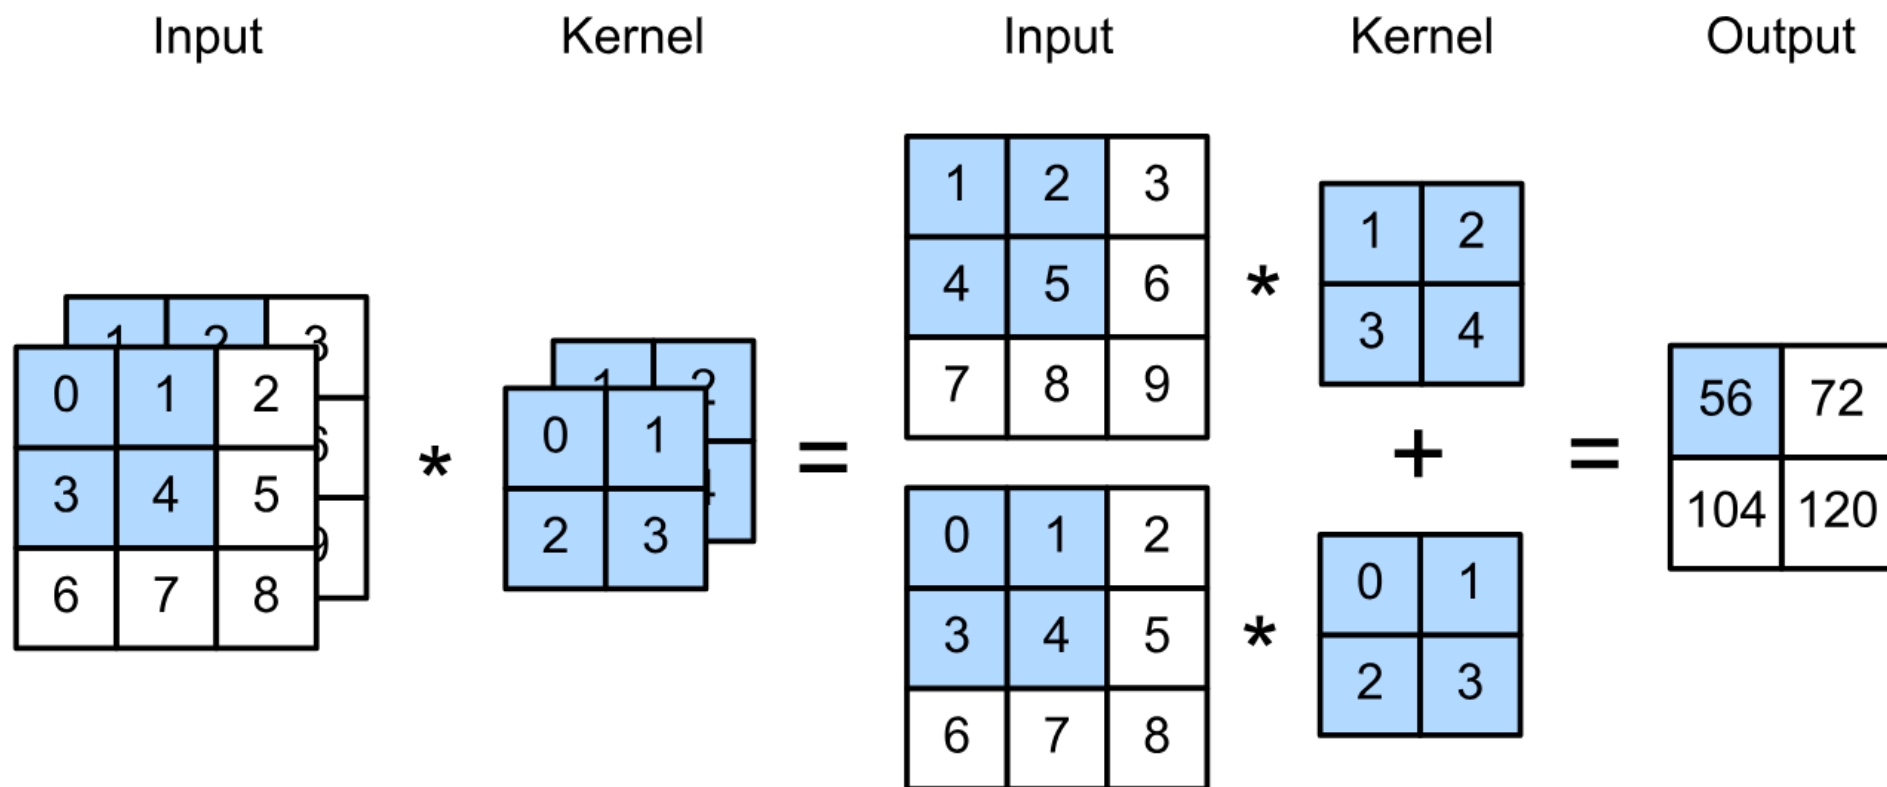


Figure 14-6. Two convolutional layers with multiple filters each (kernels), processing a color image with three color channels; each convolutional layer outputs one feature map per filter

# Two Input Channel Examples


20



# $1 \times 1$ Convolutional Layer

21

- With a  $1 \times 1$  convolution, the convolution is only accessing one pixel, and no adjacent pixels
- But this one pixel consists of multiple channels
- $\Rightarrow$  like performing a dot-product of channel vector with kernel's vector



# Notebook


22

□ `chapter_convolutional-neural-networks/channels.ipynb`

# Pooling

23

- Pooling combines multiple pixels into one pixel, while keeping the number channels the same size
- Rational for pooling
  - ▣ Down sample to reduce size
  - ▣ Mitigating sensitivity of convolution to specific location
- Operations to combine pixels: **max**, average, min



# Pooling and Stride

24

- The default stride for a pooling is the same shape as the pooling window shape
- E.g., if the pooling window shape is  $2 \times 2$ , then the stride is  $2 \times 2$

See `chapter_convolutional-neural-networks/pooling.ipynb`



# LeNet

25

- Developed by Yann LeCun, LeNet is one of the first CNNs
- Designed for digit recognition, it reached error rate of 1%
- Error rate comparable to support vector machines
- LeNet is in use in ATM machine

## LeNet-5

