

DSCI 565: ATTENTION MECHANISMS AND TRANSFORMERS

*This content is protected and may not
be shared, uploaded, or distributed.*

Ke-Thia Yao

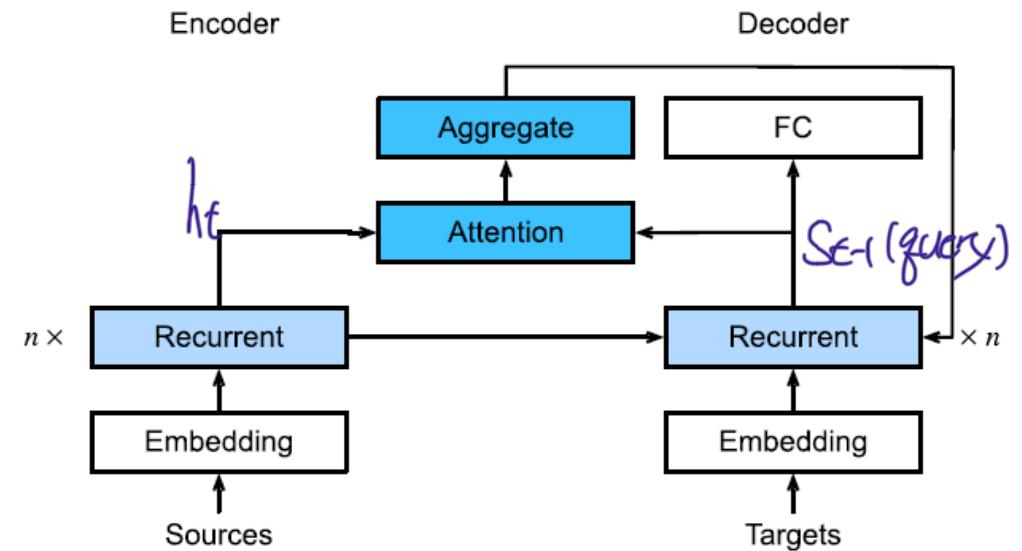
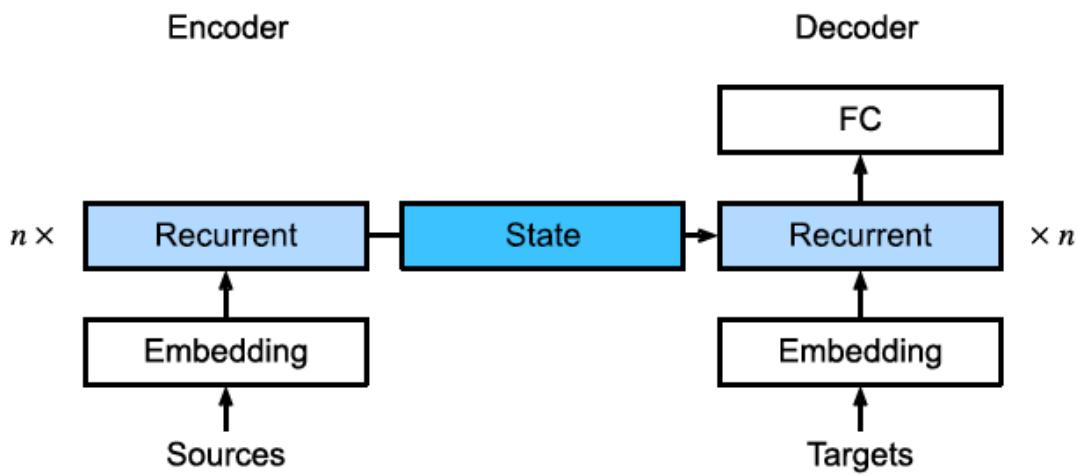
Lecture 15: 2025 October 15

Bahdanau Attention Mechanism

2

- Encoder context is no longer static (e.g., last state of the encoder)
- Context is dependent on the state $s_{t'}$ of the decoder
- Context is a weighted average of the encoder hidden states

$$c_{t'} = \sum_{t=1}^T \alpha(s_{t'-1}, h_t) h_t$$



Notebook

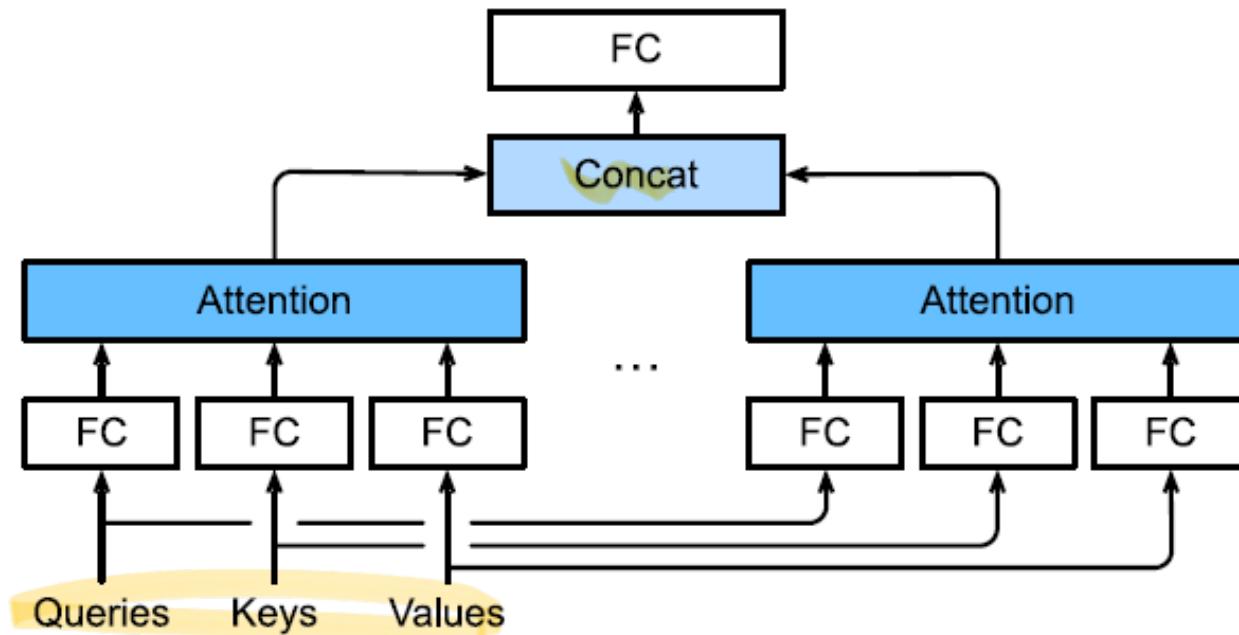
3

- chapter_attention-mechanisms-and-transformers/bahdanau-attention.ipynb

Multi-Head Attention

4

- Instead of just one attention, we may want multiple attentions
- With each attention focusing on various behavior aspects, e.g.
 - Shorter range vs. longer range
 - Parts of speech vs word meaning
- Use fully-connected layer to learn which aspects to focus using projections of queries, keys and values



Multi-Head Attention

5

- Multi-head attention transform the queries, keys and values using learned linear projections
- These *projections* are feed into *h attention poolings* in parallel
- Then the *h outputs* (\mathbf{h}_i) are concatenated
- These outputs are called *heads*
- The concatenate outputs are feed into a *fully-connected layer*

- Given a query \mathbf{q} , key \mathbf{k} and value \mathbf{v} , compute each head:

$$\mathbf{h}_i = f(\mathbf{W}_i^{(q)} \mathbf{q}, \mathbf{W}_i^{(k)} \mathbf{k}, \mathbf{W}_i^{(v)} \mathbf{v}) \in \mathbb{R}^{p_v},$$

- Combine

$$\mathbf{W}_o \begin{bmatrix} \mathbf{h}_1 \\ \vdots \\ \mathbf{h}_h \end{bmatrix} \in \mathbb{R}^{p_o}.$$

- To avoid extra computational cost, typically $p_v = \frac{p_o}{h}$

Implementation

6

```
class MultiHeadAttention(d2l.Module): #@save
    """Multi-head attention."""
    def __init__(self, num_hiddens, num_heads, dropout, bias=False, **kwargs):
        super().__init__()
        self.num_heads = num_heads
        self.attention = d2l.DotProductAttention(dropout)
        self.W_q = nn.LazyLinear(num_hiddens, bias=bias)
        self.W_k = nn.LazyLinear(num_hiddens, bias=bias)
        self.W_v = nn.LazyLinear(num_hiddens, bias=bias)
        self.W_o = nn.LazyLinear(num_hiddens, bias=bias)
```

On linear project for all heads

```
def forward(self, queries, keys, values, valid_lens):
    # Shape of queries, keys, or values:
    # (batch_size, no. of queries or key-value pairs, num_hidden)
    # Shape of valid_lens: (batch_size,) or (batch_size, no. of queries)
    # After transposing, shape of output queries, keys, or values:
    # (batch_size * num_heads, no. of queries or key-value pairs,
    # num_hidden / num_heads)
    queries = self.transpose_qkv(self.W_q(queries))
    keys = self.transpose_qkv(self.W_k(keys))
    values = self.transpose_qkv(self.W_v(values))
```

Shape: b, q, d

Map and transpose query.
Shape: $b, q, d \rightarrow b * h, q, d/h$

Similar transformation for keys
and values

```
if valid_lens is not None:
    # On axis 0, copy the first item (scalar or vector) for num_heads
    # times, then copy the next item, and so on
    valid_lens = torch.repeat_interleave(
        valid_lens, repeats=self.num_heads, dim=0)
```

Shape of output: (batch_size * num_heads, no. of queries,
num_hidden / num_heads)

```
output = self.attention(queries, keys, values, valid_lens)
```

Shape of output_concat: (batch_size, no. of queries, num_hidden)

```
output_concat = self.transpose_output(output)
```

```
return self.W_o(output_concat)
```

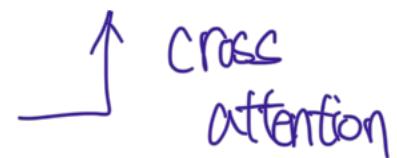
Compute in parallel:
all batches and all heads

Reverse transpose_qkv

Notebook

8

- chapter_attention-mechanisms-and-transformers/multihead-attention.ipynb



Cross
attention

Self-Attention

9

- Given sequence $\mathbf{x}_1, \dots, \mathbf{x}_n$, self-attention outputs a sequence $\mathbf{y}_1, \dots, \mathbf{y}_n$

$$\mathbf{y}_i = \text{Attention}(\mathbf{x}_i, \mathcal{D})$$

where $\mathcal{D} = \{(\mathbf{x}_1, \mathbf{x}_1), \dots, (\mathbf{x}_n, \mathbf{x}_n)\}$

- The queries, keys and values are all \mathbf{x}_i
- The idea is to find other parts of the sequence $\mathbf{x}_1, \dots, \mathbf{x}_n$ to which to pay attention

Self-Attention

10

For self-attention use
the same matrix for all
three inputs Q, K, V

For example, use the
Input sentence
embedding matrix for
all three inputs

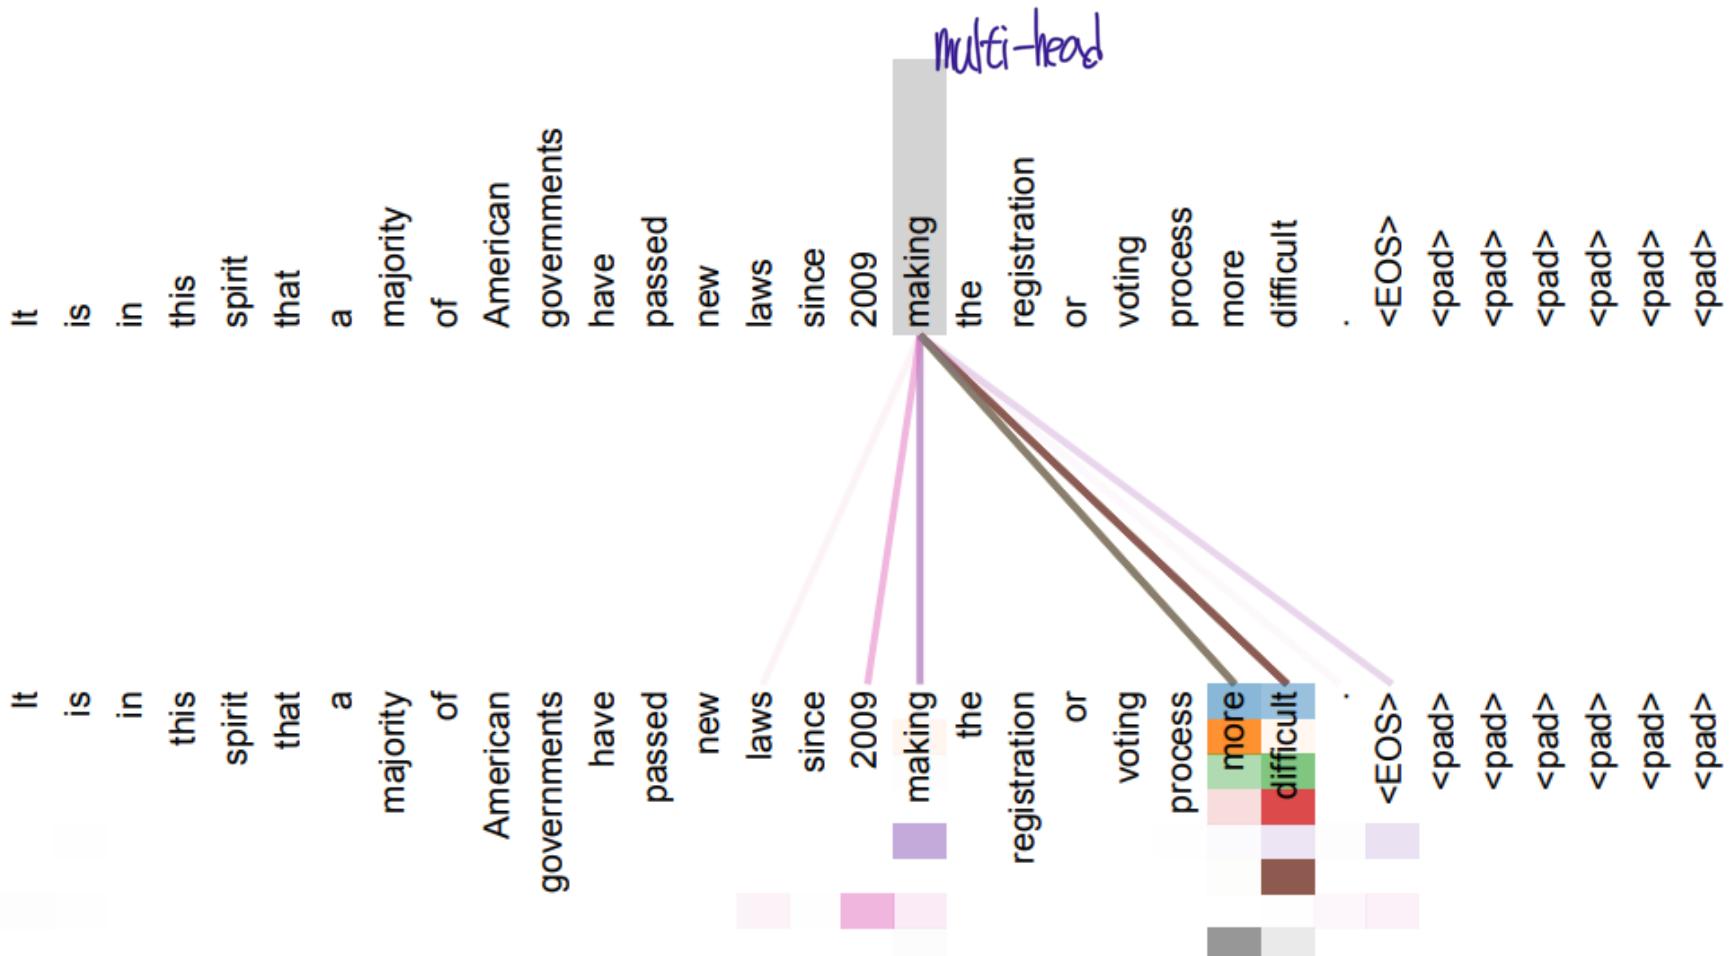


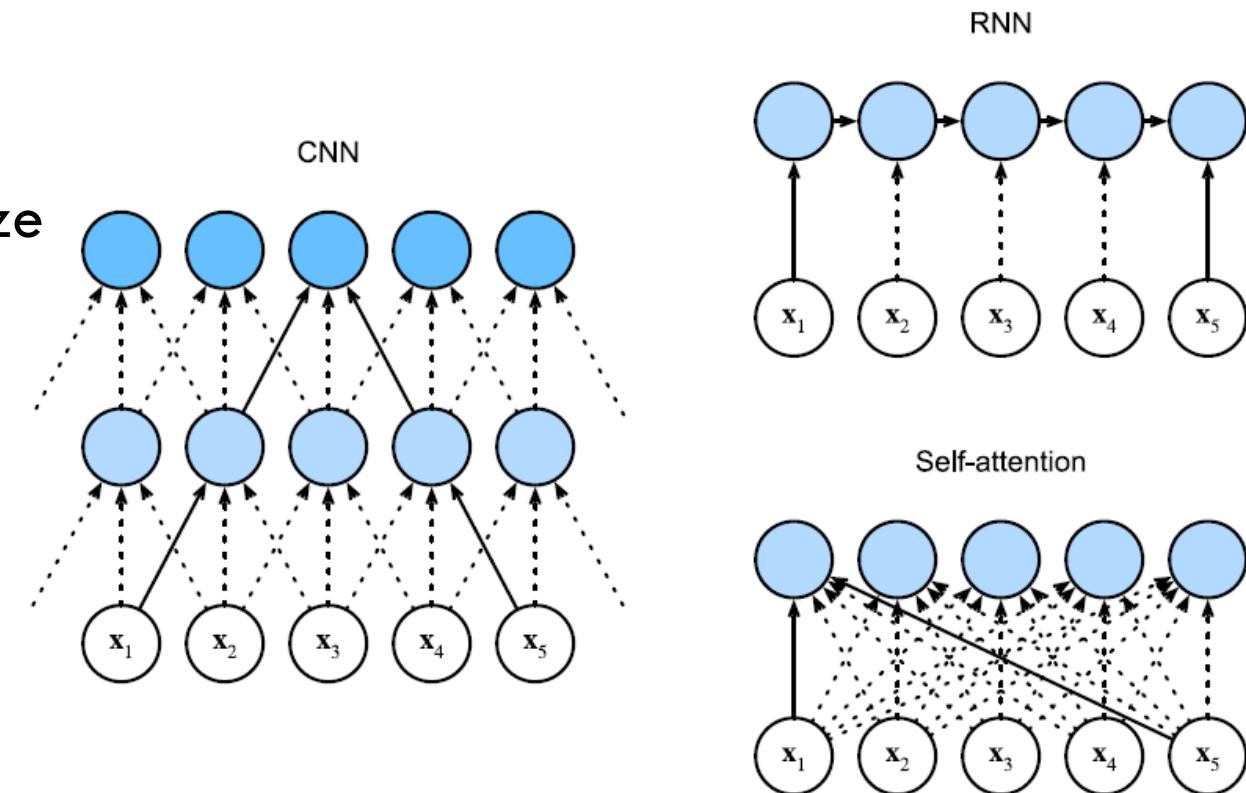
Figure 3: An example of the attention mechanism following long-distance dependencies in the encoder self-attention in layer 5 of 6. Many of the attention heads attend to a distant dependency of the verb ‘making’, completing the phrase ‘making...more difficult’. Attentions here shown only for the word ‘making’. Different colors represent different heads. Best viewed in color.

Comparing CNNs, RNNs and Self-Attention

11

- All can be thought as ways to propagate information along a sequence
- Let n be the sequence length
 d be the input/output channel, hidden size
 k be the kernel size

	CNN	RNN	Self-Attention
Complexity (per layer)	$\mathcal{O}(knd^2)$	$\mathcal{O}(nd^2)$	$\mathcal{O}(n^2d)$
Sequential Ops	$\mathcal{O}(1)$	$\mathcal{O}(n)$	$\mathcal{O}(1)$
Max path length	$\mathcal{O}(n/k)$	$\mathcal{O}(n)$	$\mathcal{O}(1)$



Positional Encoding

12

- Attention does not preserve order information
- The vector y_i is just a weighted average of other vectors
- What if order information is needed for the translation task?
- Could just concatenate an index (in binary) to each word embedding
- Positional encoding does something similar, except
 - It uses sine and cosine functions
 - Adds index to embedding, instead of concatenation

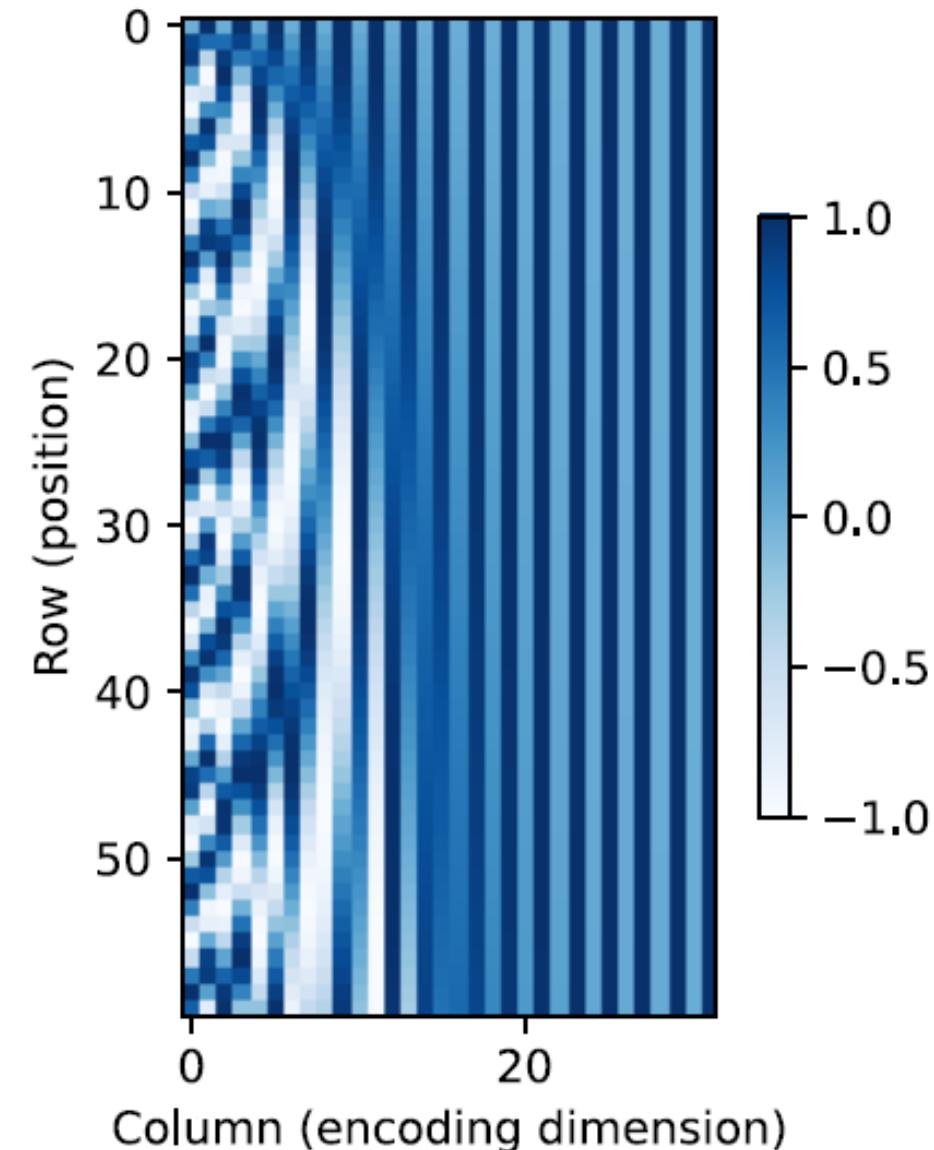
0000
0001
0010
0011
0100
0101
0110
0111
1000
1001
1010
1011
1101
1110
1111

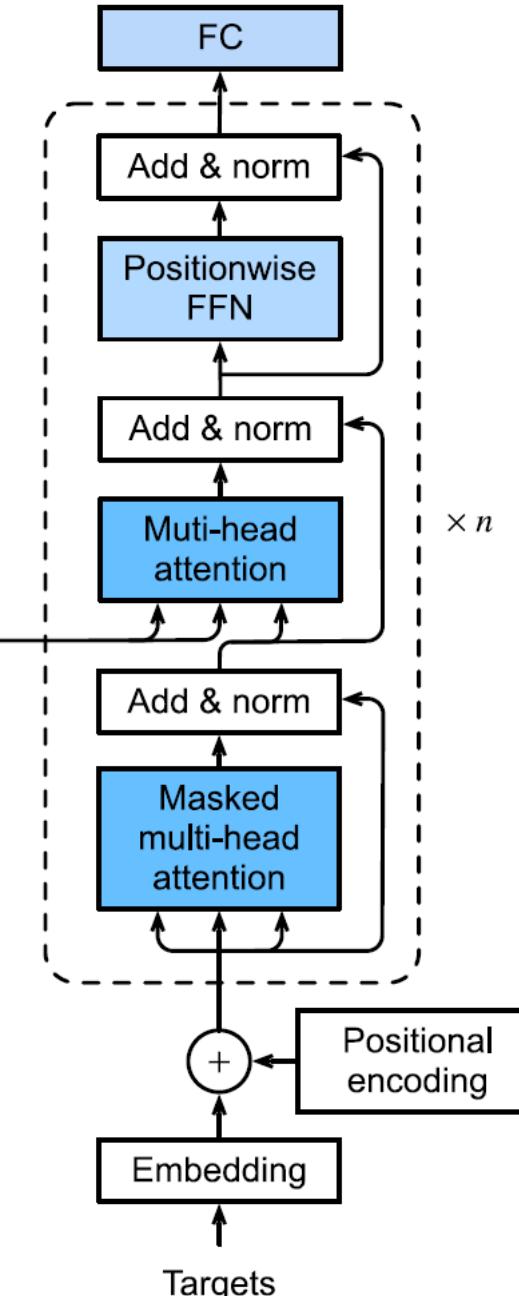
Positional Encoding

13

- Let $\mathbf{X} \in \mathbb{R}^{n \times d}$ be the input sequence and $\mathbf{P} \in \mathbb{R}^{n \times d}$ be in position encoding, the positional encoded output is $\mathbf{X} + \mathbf{P}$
- Where $i \in [0, n)$ and $j \in [0, \frac{d}{2}]$

$$p_{i,2j} = \sin\left(\frac{i}{10000^{2j/d}}\right)$$
$$p_{i,2j+1} = \cos\left(\frac{i}{10000^{2j/d}}\right)$$





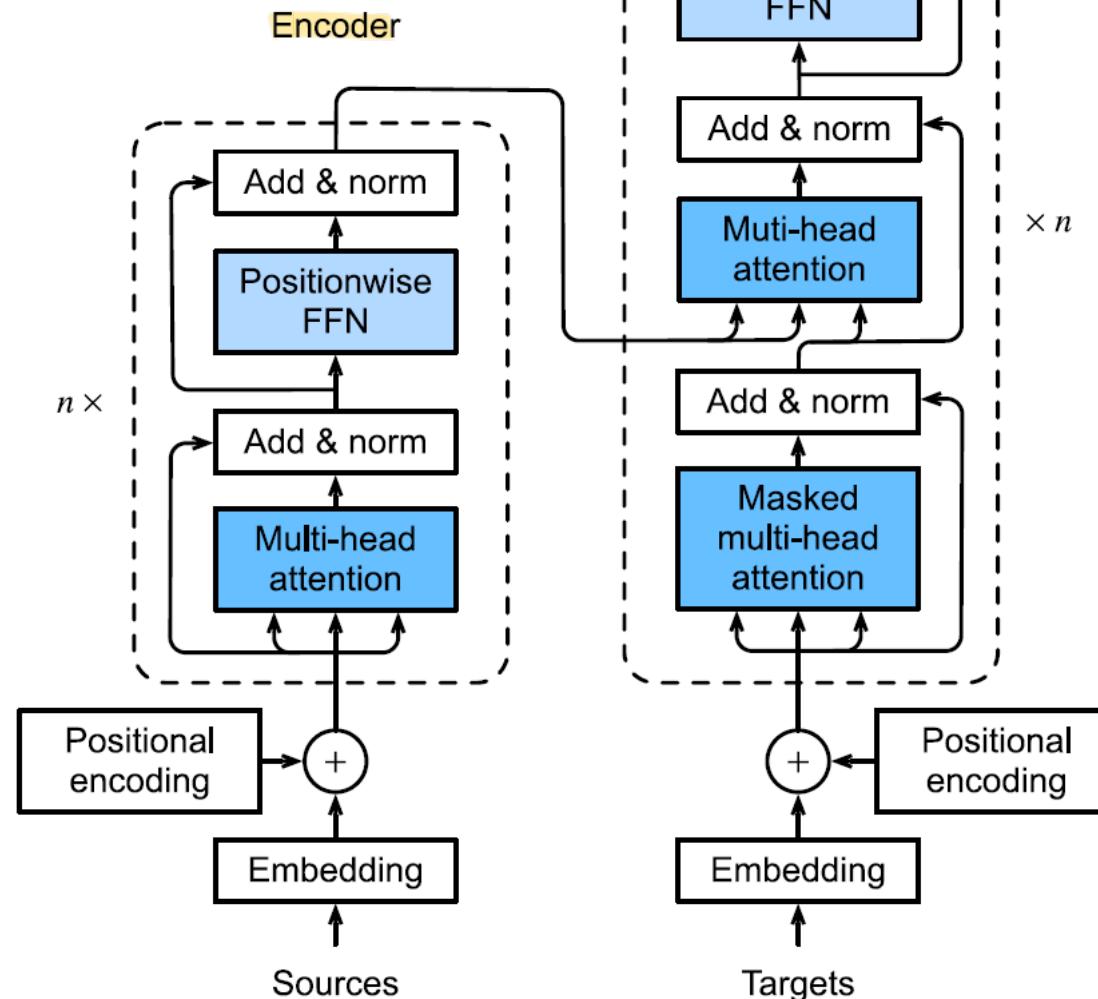
Transformer Architecture

- “Attention is All You Need,” Vaswani et al., 2017
- Regular neural network with attention mechanism
- No recurrent units
 - Easier to train
 - No vanishing gradients
 - Can be parallelized
- Pervasive in deep learning applications: language, vision, speech and reinforcement learning

Training and Prediction

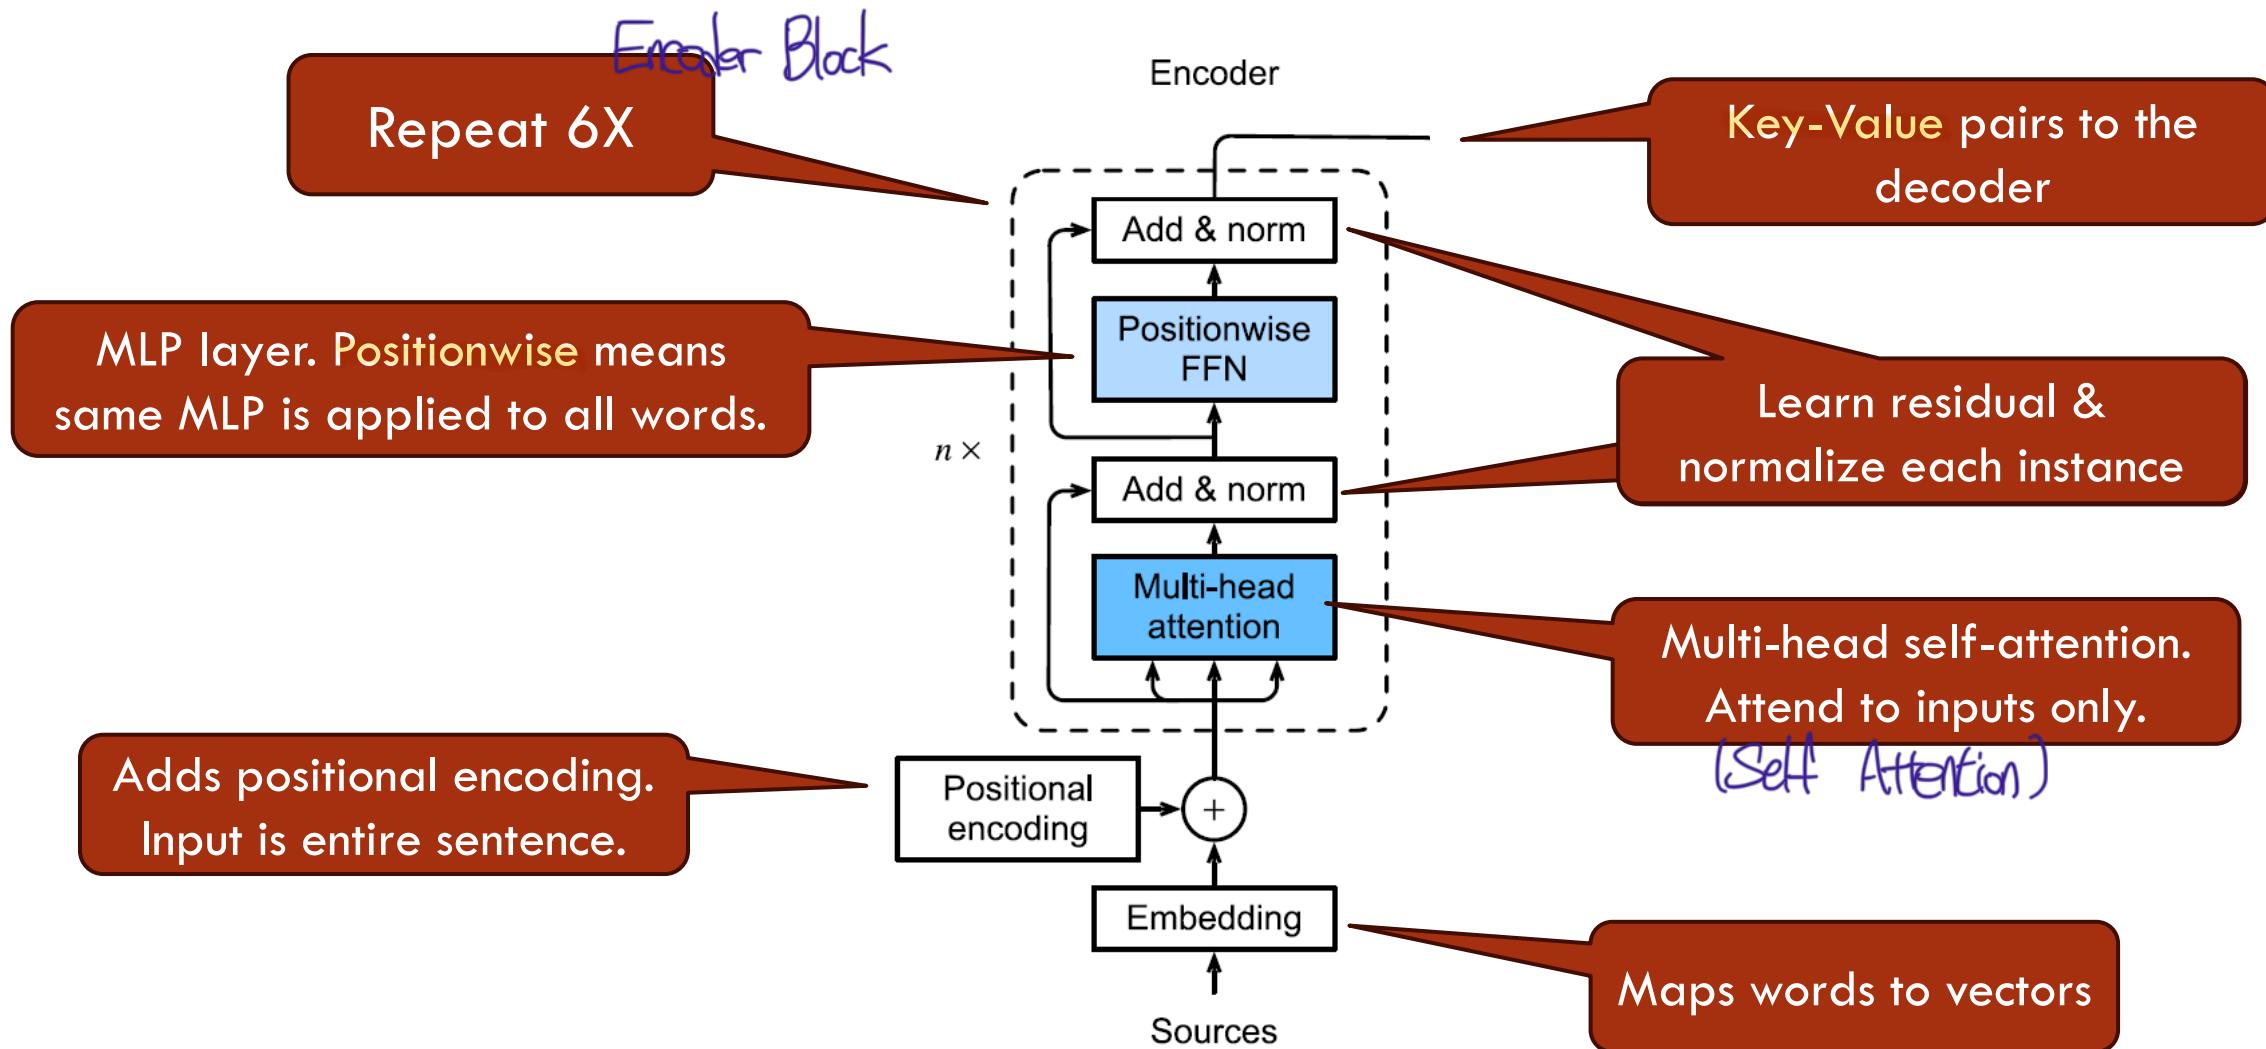
- Training phase
 - Send entire English sentence to Encoder, and Spanish sentence to decoder
 - “I like soccer” and “<sos> Me gusta el futbol”

- During translation phase *Interence phase*
 - Call transformer multiple times
 - “I like soccer” and “<sos>”
 - “I like soccer” and “<sos> Me”
 - “I like soccer” and “<sos> Me gusta”
 - “I like soccer” and “<sos> Me gusta el”



Encoder

16



Decoder

17

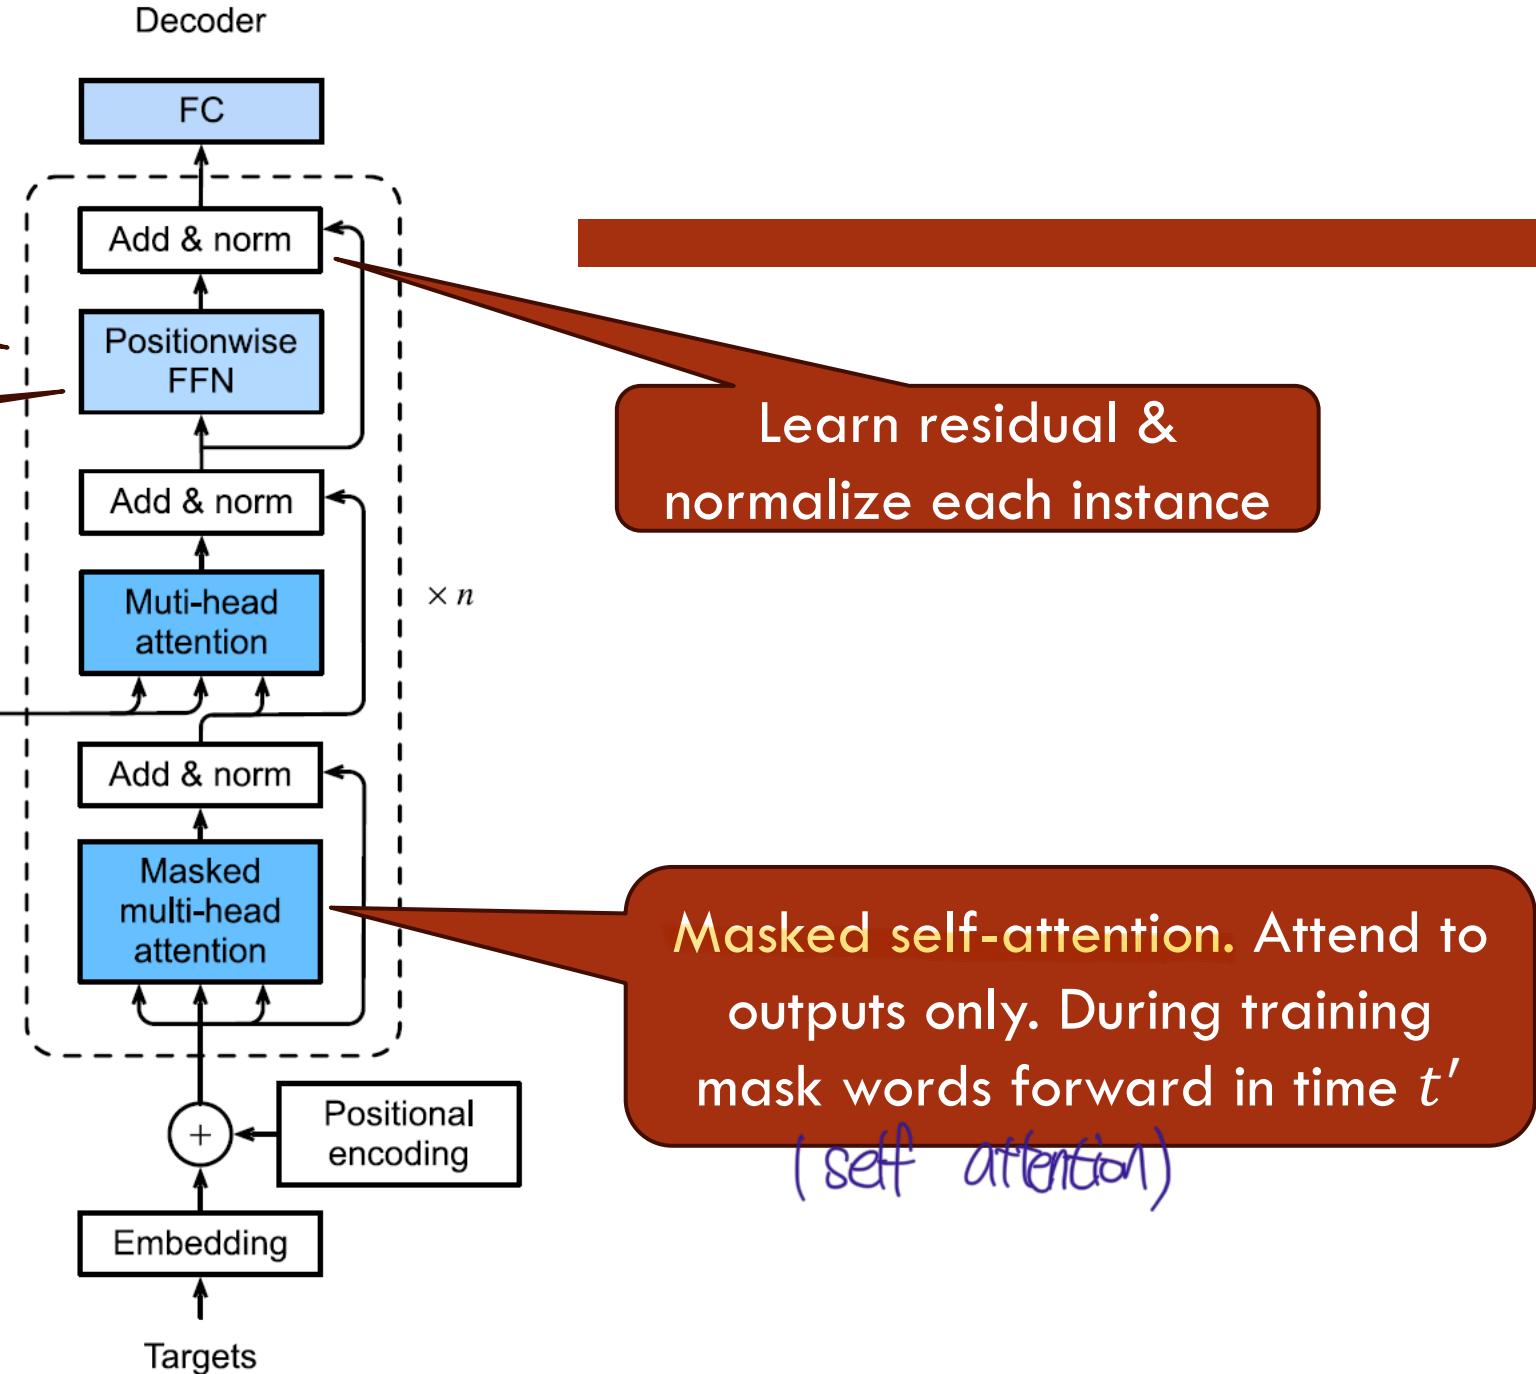
Repeat 6X

Decoder Block

MLP layers. Positionwise means same MLP is applied to all words.

Cross attention. Use decoder states to attend to encoder states.
(cross attention)

key, value from
encoder



Notebook

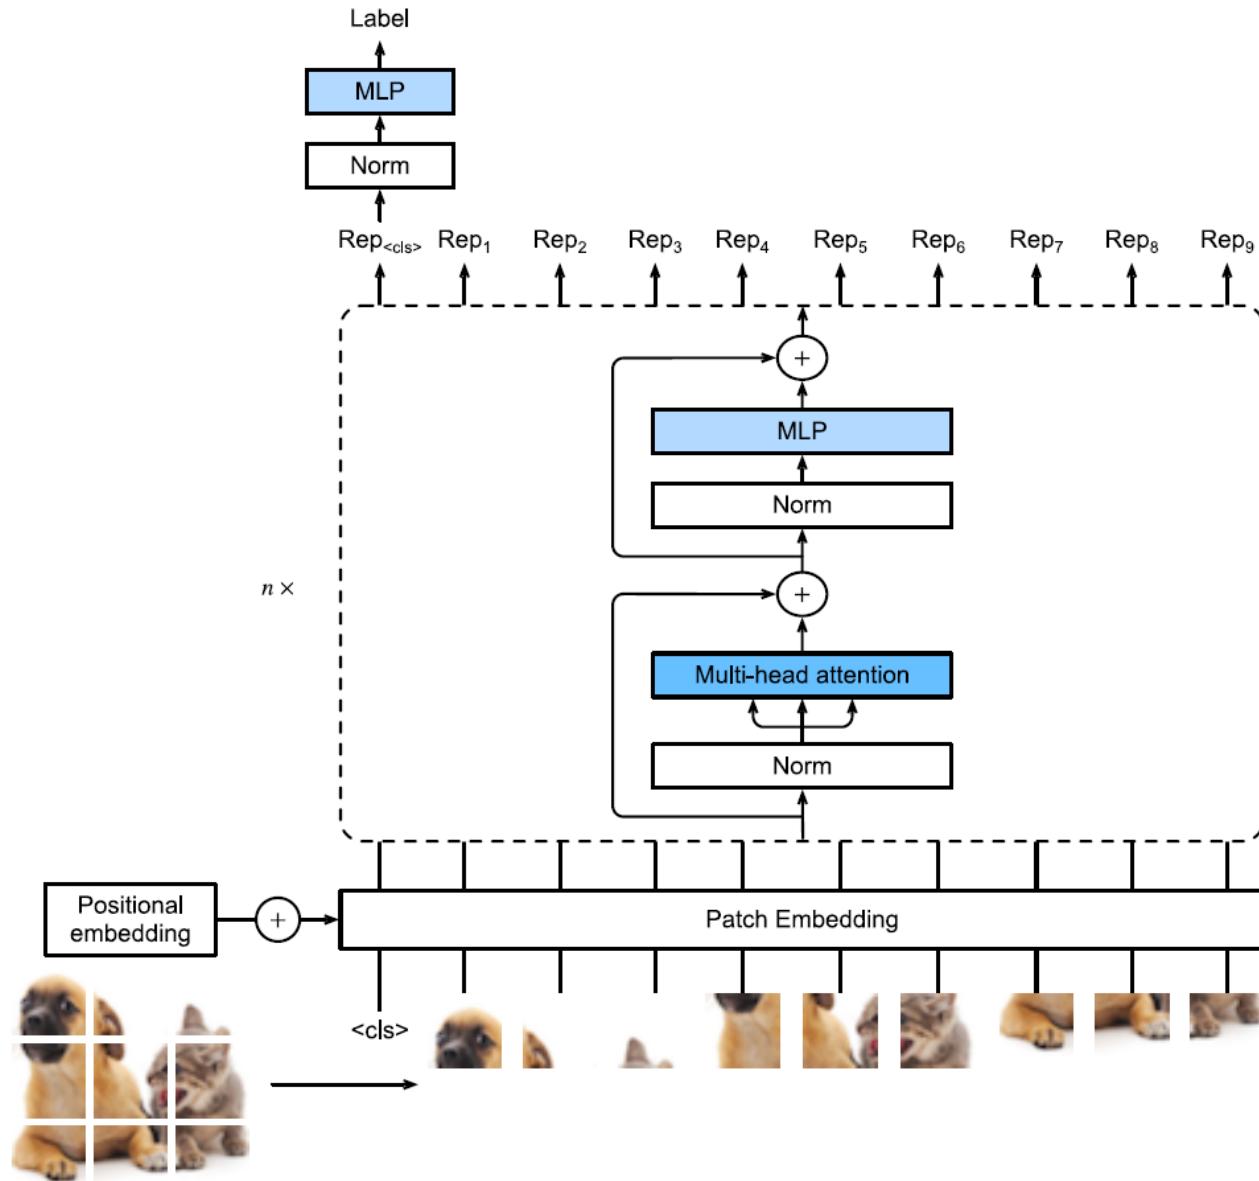
18

- chapter_attention-mechanisms-and-transformers/transformer.ipynb

Vision Transformers

19

- Can the transformer architecture be applied to computer vision?
- Vision Transformers (ViTs), Dosovitskiy et al., 2021
 - Better scalability than CNN
 - Performs better than ResNets on larger datasets

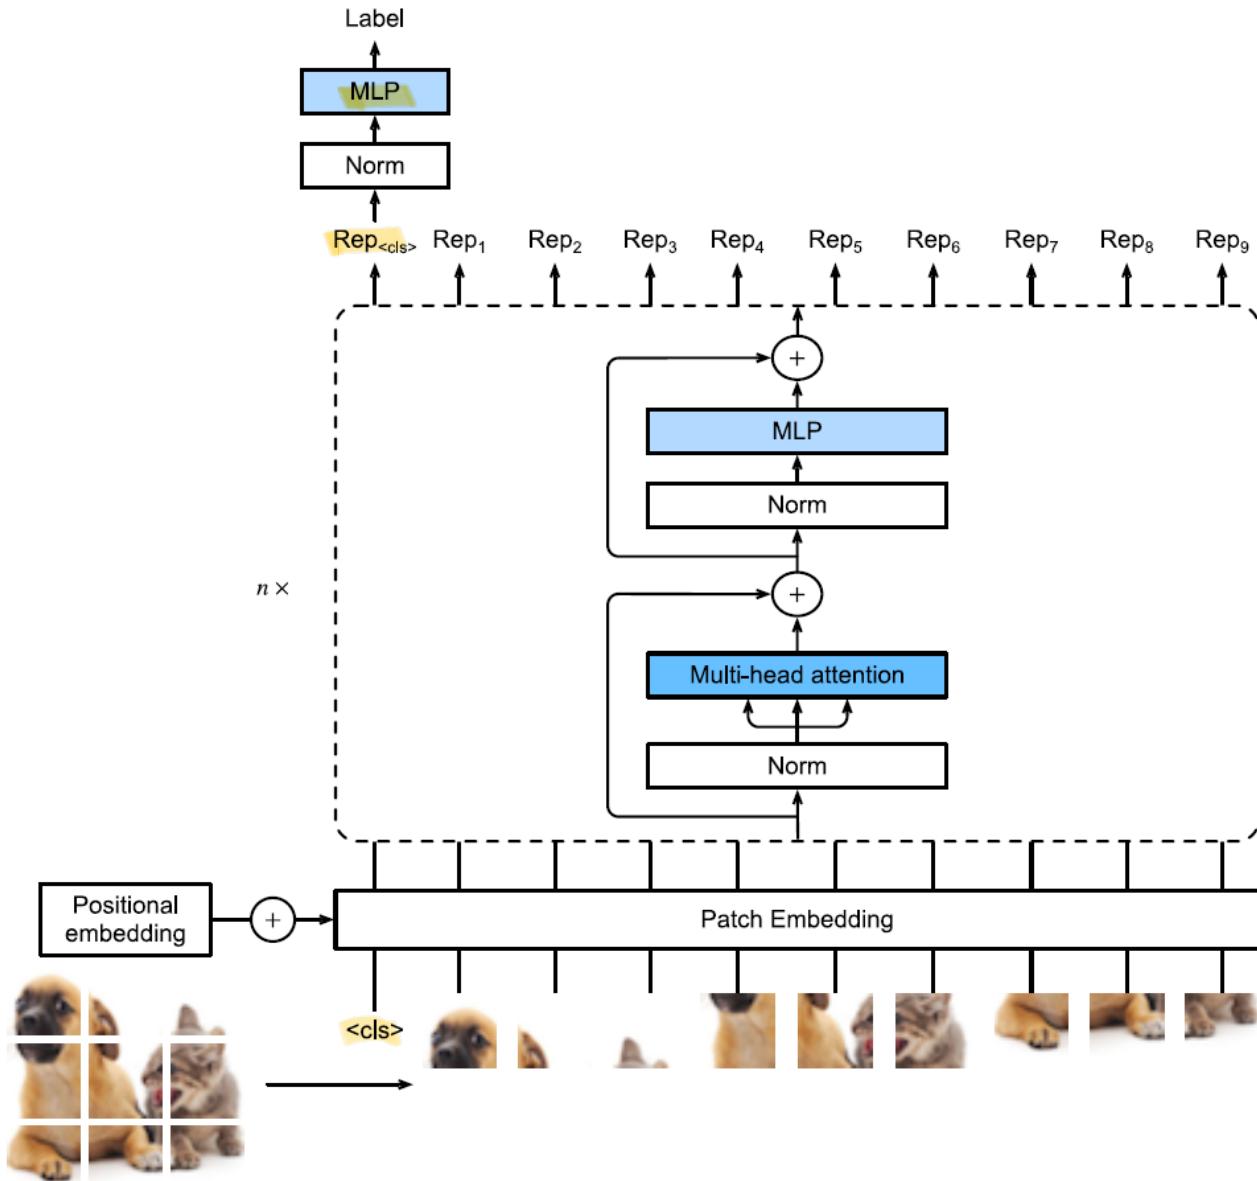


Vision Transformers

only encoder

20

- Partition image into non-overlapping patches
 - Conv with kernel size = stride
- Embed patches as vectors
- Add horizontal and vertical positional embedding
- Self-attention across all patches
- Layer normalization with learning residuals
- <cls> token with corresponding $\text{Rep}_{<\text{cls}>}$ for classification



Improving Vision Transformers

21

- Improving ViT on smaller datasets
 - DeiT (Touvron *et al.*, 2021) more data efficient training strategies
 - T2T ViT (Yuan *et al.*, 2021) use token-to-token mappings
- Scaling ViT to larger images
 - Swin Transformers (Liu *et al.*, 2021) hierarchical transformers with shifted windows to avoid global self-attention

Notebook

22

- chapter_attention-mechanisms-and-transformers/vision-transformer.ipynb

Large-Scale Pretraining with Transformers

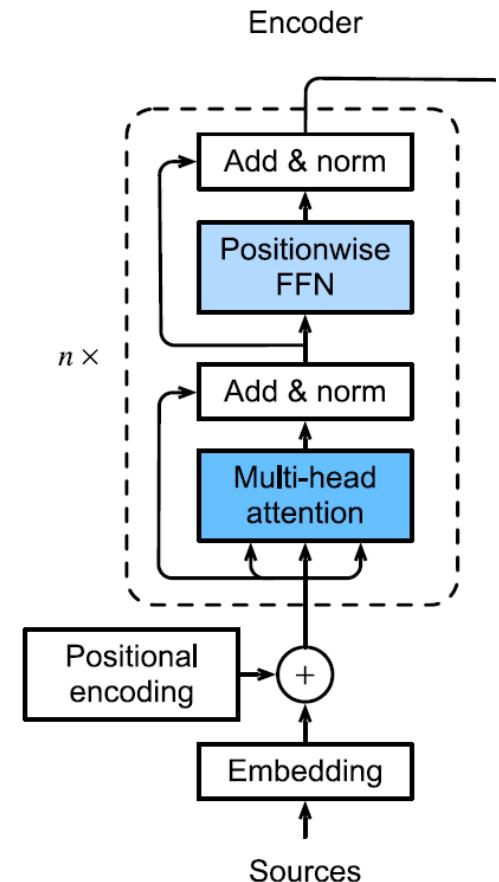
23

- To build a machine learning model for a specific tasks, we could train the model *from scratch* using from a task-specific dataset
- A potential problem of the from-scratch approach is the resulting model may be sensitive to shifts in the data distribution
 - overfitting*
- Another approach is to first train a *general model* on a large dataset. Then, use *transfer learning* to adapted the *pretrained model* to the specific task
 - overfitting*
- This second approach is used by researcher with CNN (e.g., EfficientNet, ResNet)
- This is the approach taken by researcher with transformers as well

Encoder-Only Pretrained Model

24

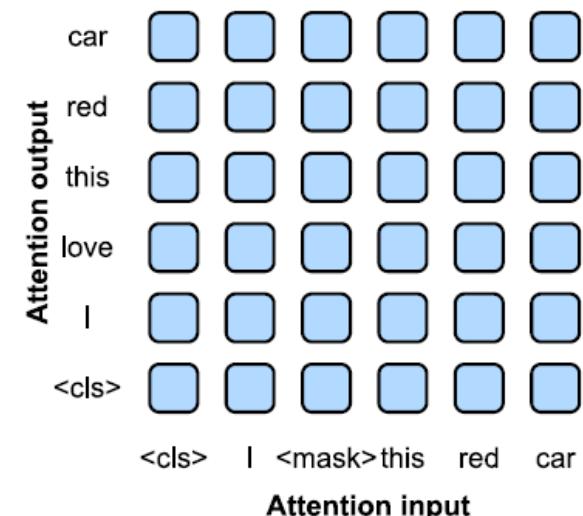
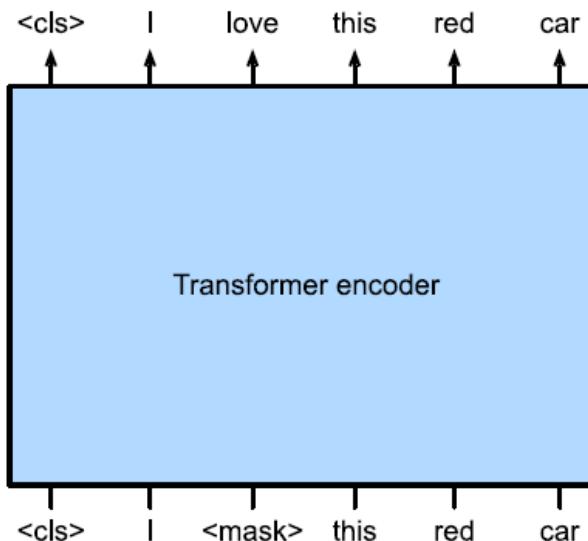
- The Transformer encoder converts a sequence of input tokens into a vector representation
- Then the representation can be used further for learning, such as generated a classifier
- The vision Transformer is an encoder-only architecture with the representation of the special token <cls> used for classification



BERT: Bidirectional Encoder Representations from Transformers

25

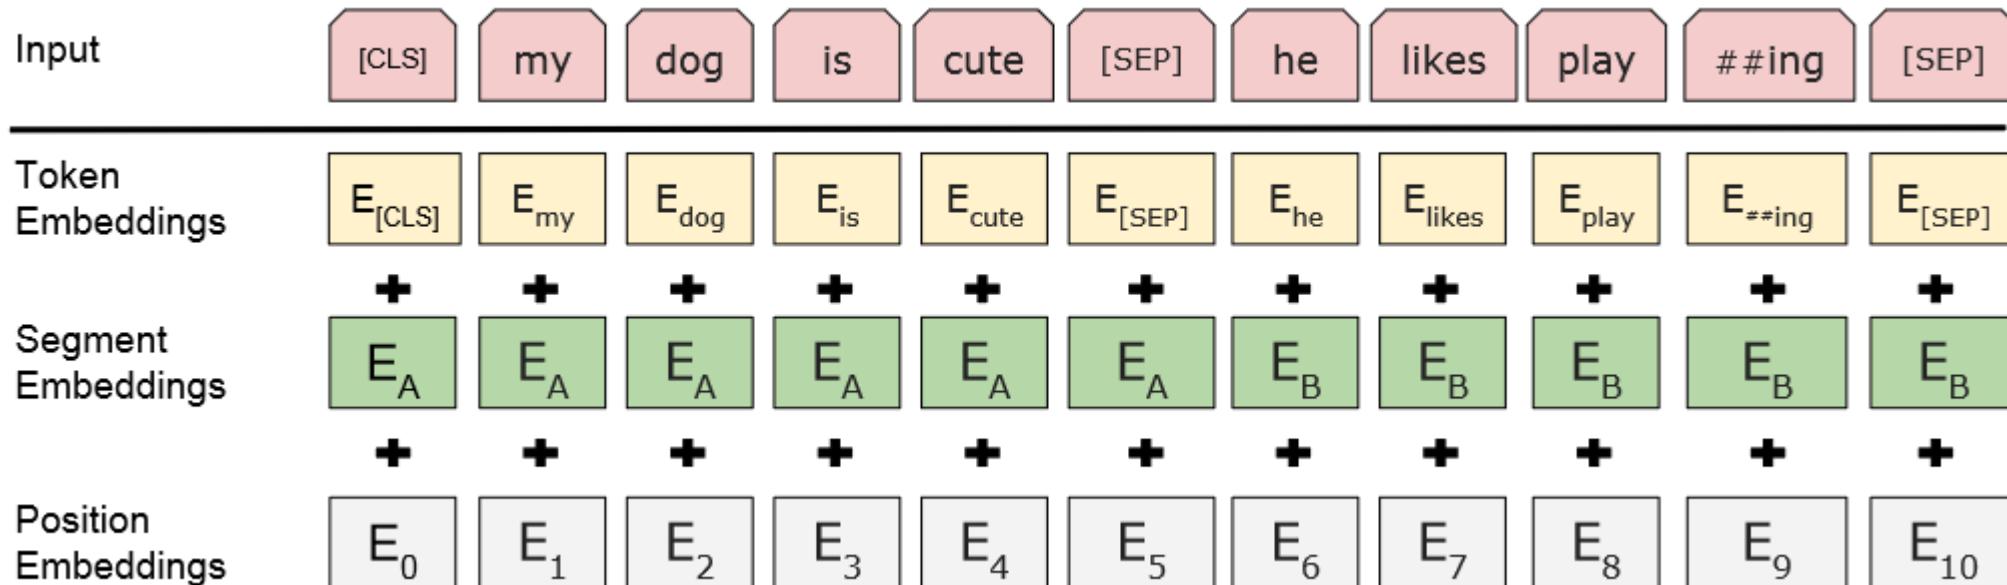
- BERT is pre-trained model on text sequences using *masked language modeling*
- BERT has 350M parameters
- BERT is pre-trained on BooksCorpus (800M words) and English Wikipedia (2,500M words)
- Pre-training Task #1: Masked LM. Tokens from the input text sequences are randomly masked
 - Masked input: <cls> I <mask> this red car
- During training BERT is given the entire masked input. This allows all tokens to attend to each other (i.e., “bidirectional”)



BERT

26

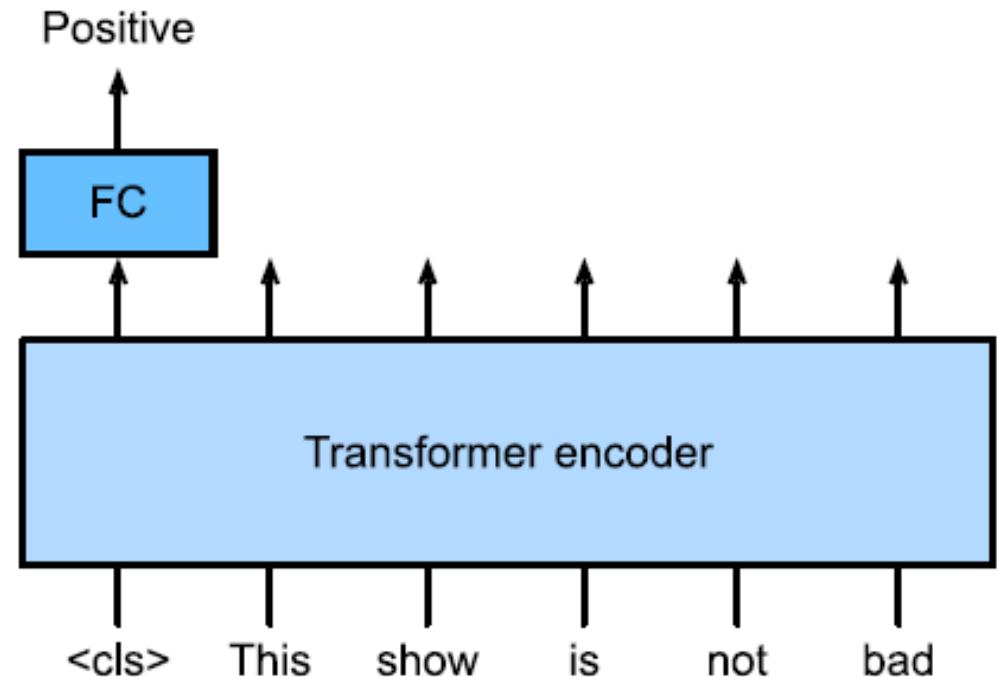
- Task #2: Next Sentence Prediction (NSP)
- Training set consists of pairs of sentences
 - 50% the second sentence is from the actual next sentence
 - 50% the second sentence is replaced by a randomly chosen sentence



Fine-Tuning BERT

27

- To fine-tune BERT, add additional layers with randomized parameters
- Then train to update the new parameters as well pretrained BERT parameters
- Right shows fine-tuning BERT for sentiment analysis



Fine Tuning BERT

28

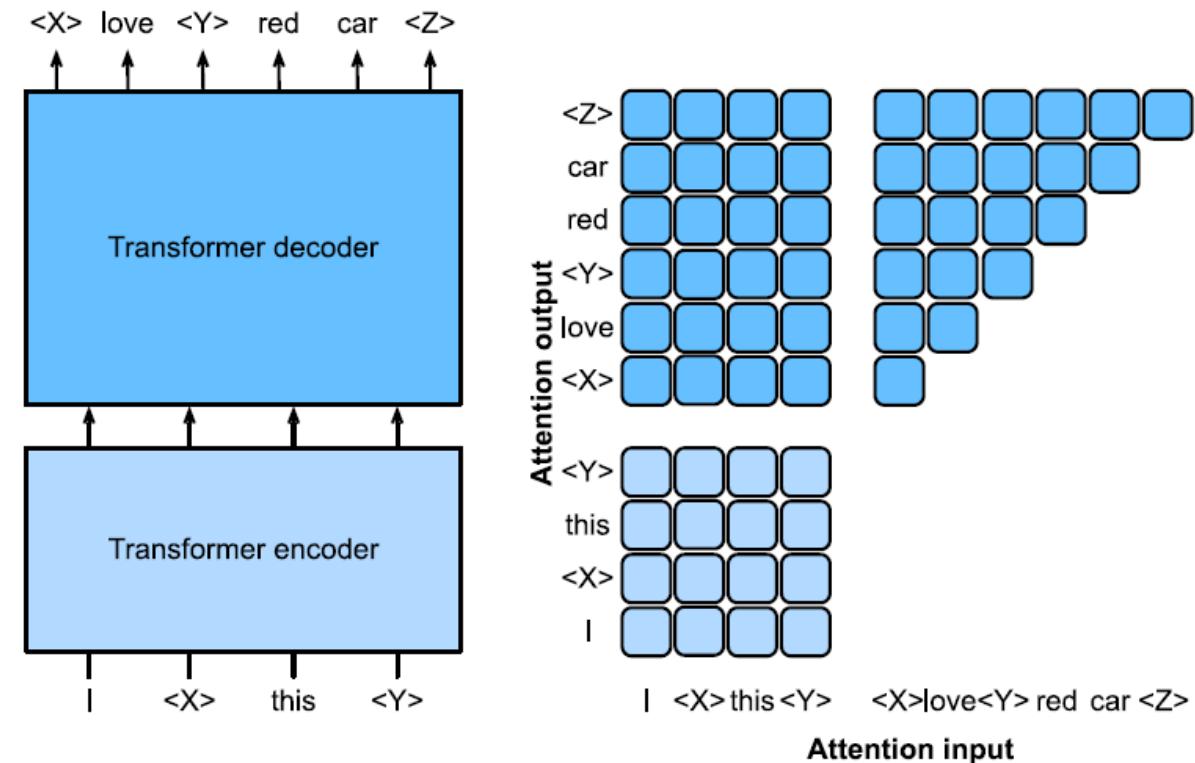
- The BERT paper (Devlin et al., 2019) fine tuned BERT on four tasks
 - GLUE benchmark for language understanding modeled as classification problem
 - SQuAD v1.1 question answering problem. Given a question and a passage from Wikipedia containing the answer, the task is to predict the answer text span in the passage.
 - SQuAD v2.0 extends v1.1 with problems with passages that do not contain the answer
 - SWAG common sense reasoning. Given sentence and choose the most plausible among four possible continuation sentences
- Required 2 to 4 epochs to fine tune for a task

mid term //

Encoder-Decoder Pretrained Model

29

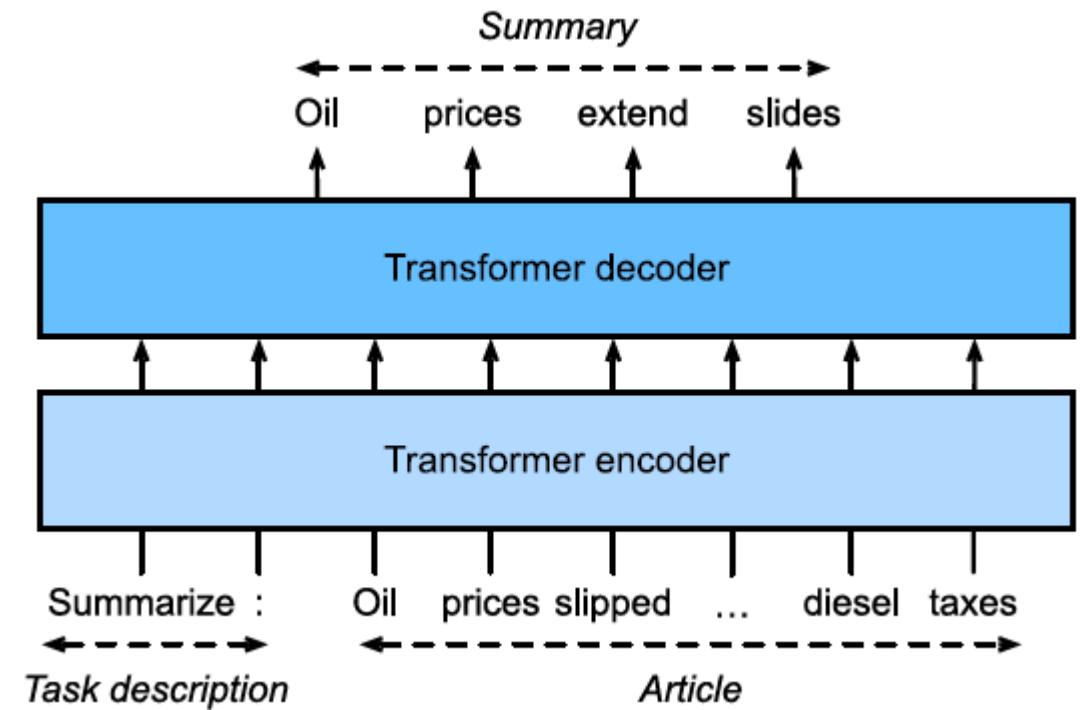
- Examples include BART (Lewis *et al.*, 2019) from Meta, and T5 (Raffel *et al.*, 2020) from Google
- T5 trained to generate missing spans of words
- Given
 - Original sentence: I love this red car
 - Input: I <X> this <Y>
 - Target: <X> love <Y> red card <Z>



Fine Tuning T5

30

- Fine tuning T5
 - Add task description to input
 - No additional layers are needed
- T5 can generate output of arbitrary length
- After fine tuning T5 (11 billion parameters), it achieved state-of-the-art performance in encoding/classification tasks and generation/summarization tasks



T5 Encoder

31



Teddy bears swimming at the Olympics 400m Butterfly event.



A cute corgi lives in a house made out of sushi.



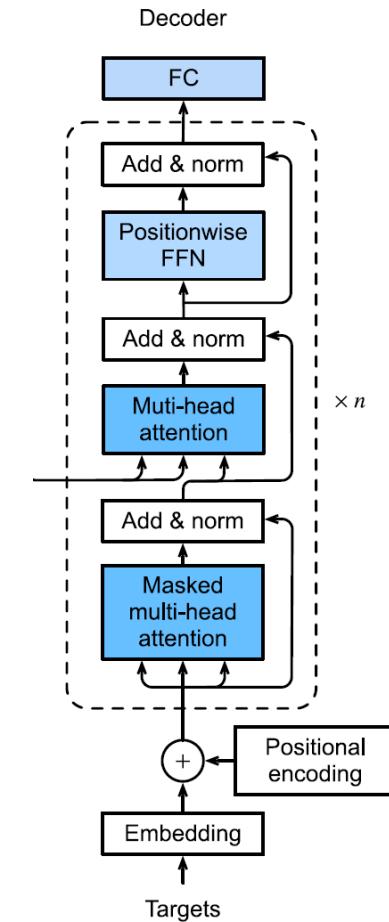
A cute sloth holding a small treasure chest. A bright golden glow is coming from the chest.

Text-to-image examples by the Imagen model, whose text encoder is from T5 (figures taken from Saharia et al. (2022)).

Decoder-Only Pretrained Models

32

- Examples include GPT (100 million parameters), GPT-2 (1.5 billion parameters) and GPT-3 (175 billion parameters)
- GPT-2 was trained on 40GB of text and obtained state-of-the-art results multiple language benchmarks *without updating the parameters or architecture*

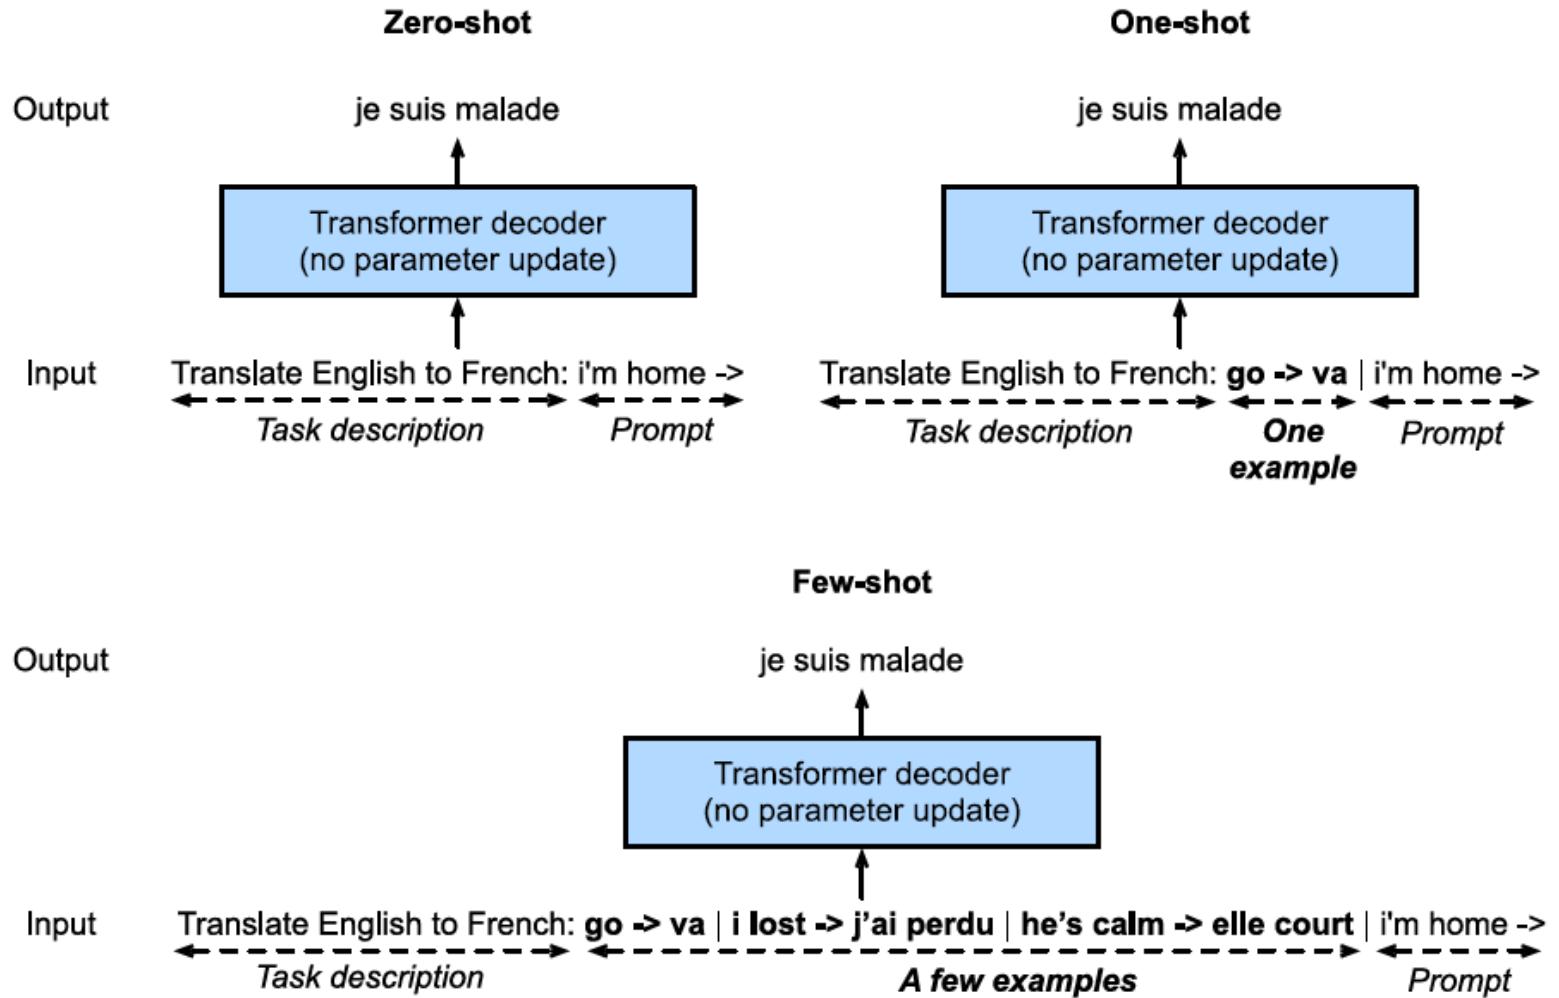


In-Context Learning

33

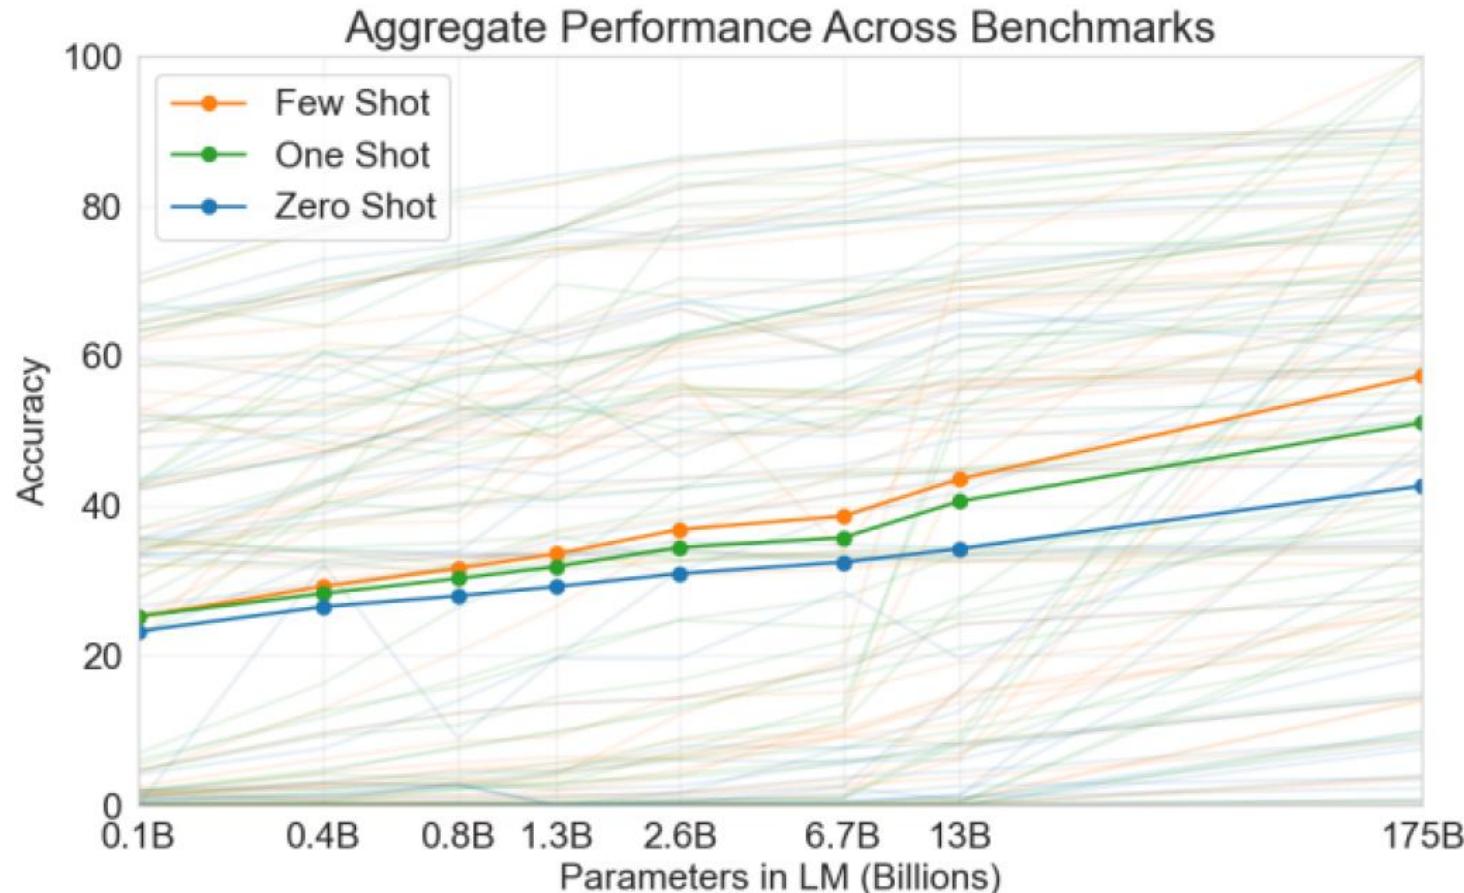
Input consist of:

- Task description
- Zero or more examples
- A Prompt



GPT-3 Performance

34

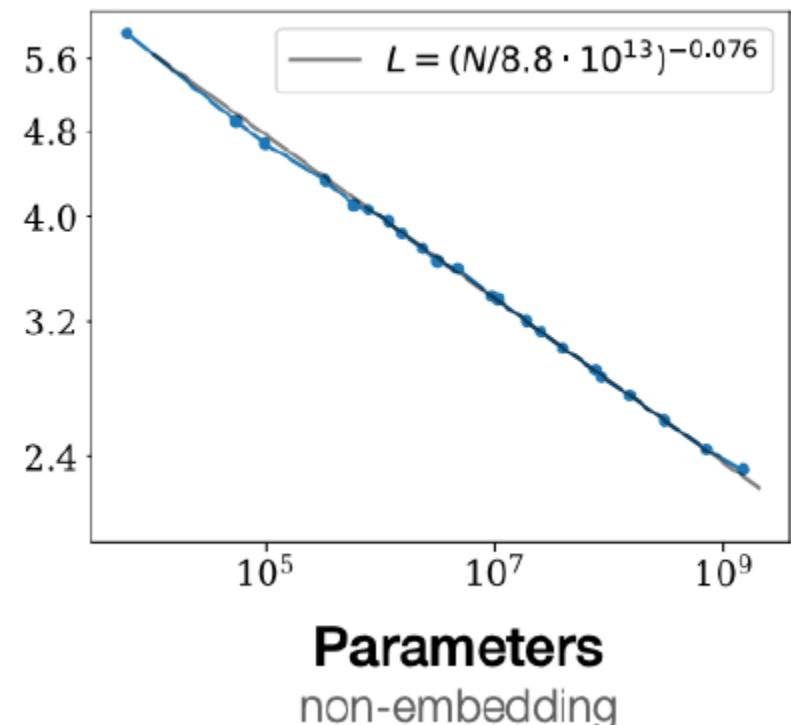
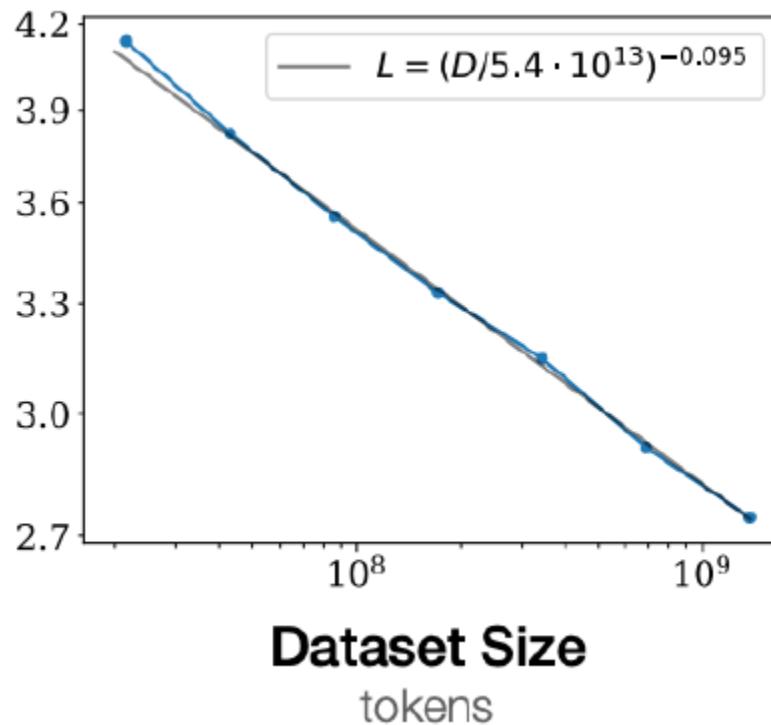
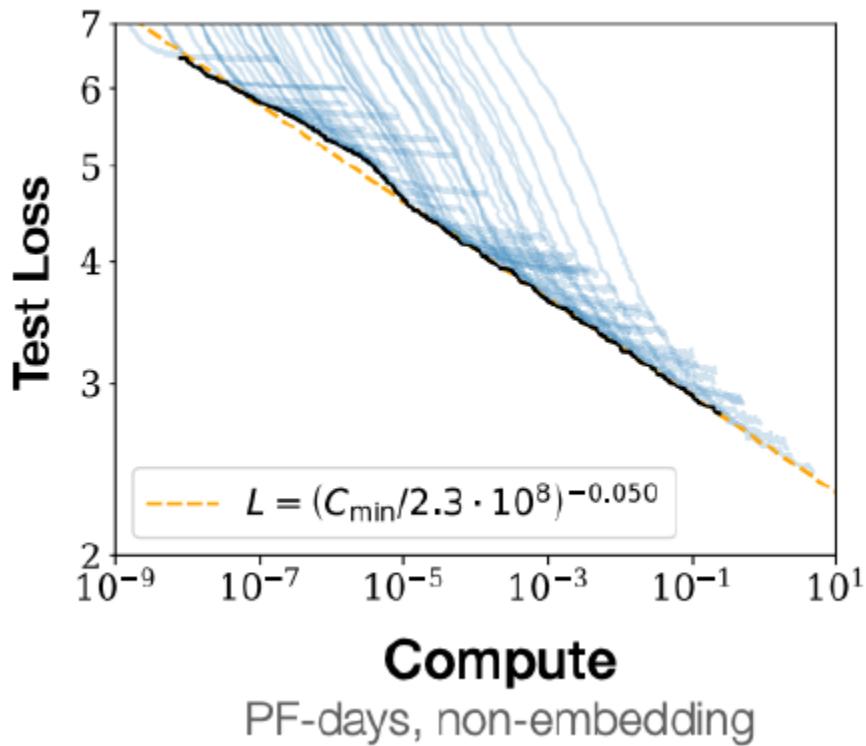


Aggregate performance of GPT-3 for all 42 accuracy-denominated benchmarks (caption adapted and figure taken from Brown et al. (2020)).

Scalability

35

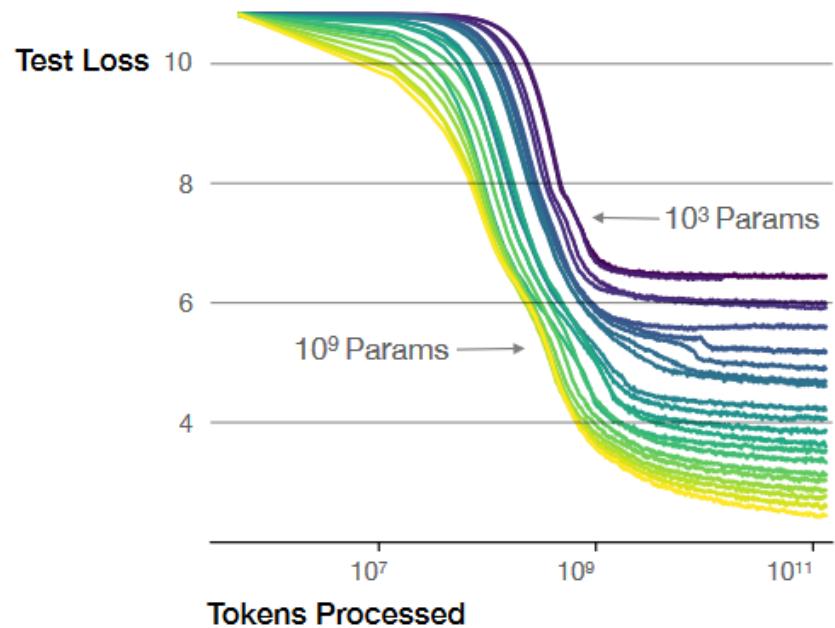
- Power-law scaling of Transformers in GPT language model (Kaplan et al., 2020)
- Test loss decreases smoothly with compute, data size and parameters



Scalability

36

Larger models require fewer samples to reach the same performance



The optimal model size grows smoothly with the loss target and compute budget

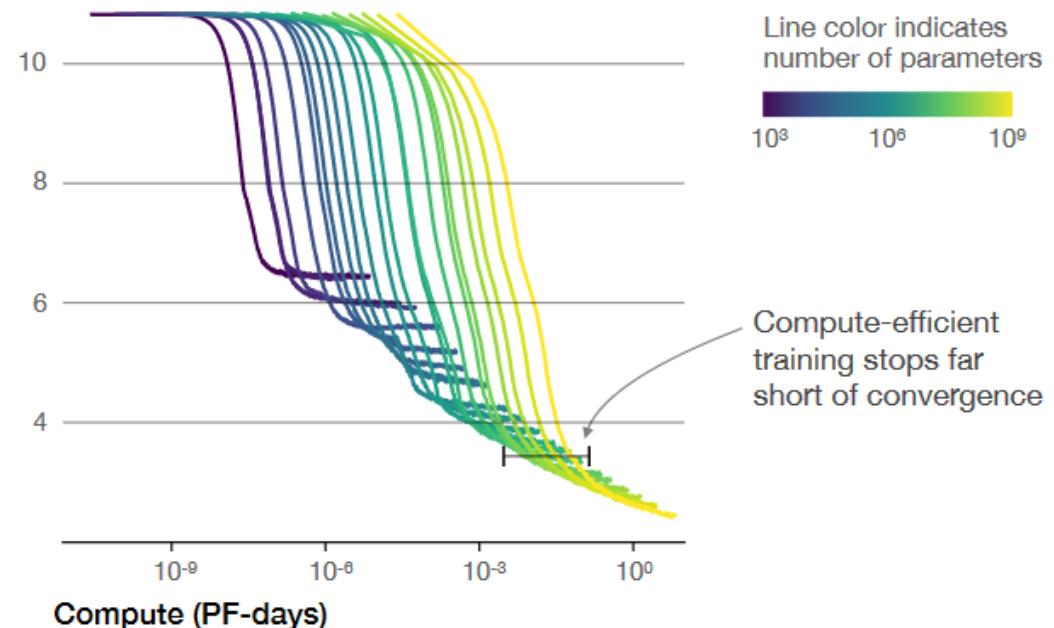


Figure 2 We show a series of language model training runs, with models ranging in size from 10^3 to 10^9 parameters (excluding embeddings).

Scalability

37

- Power law scaling persist for GTP-3 sized networks

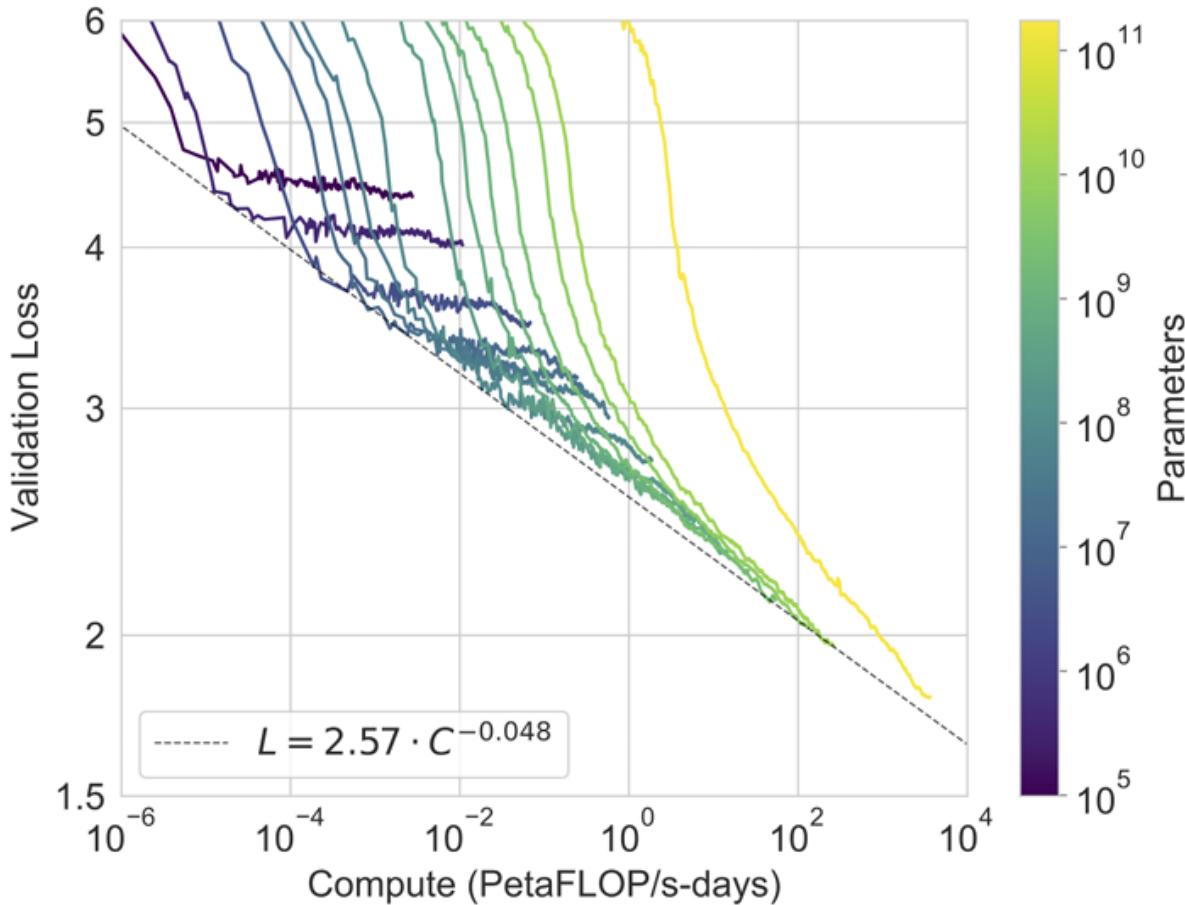


Figure 3.1: Smooth scaling of performance with compute. Performance (measured in terms of cross-entropy validation loss) follows a power-law trend with the amount of compute used for training. The power-law behavior observed in [KMH⁺20] continues for an additional two orders of magnitude with only small deviations from the predicted curve. For this figure, we exclude embedding parameters from compute and parameter counts.

Large Language Models

38

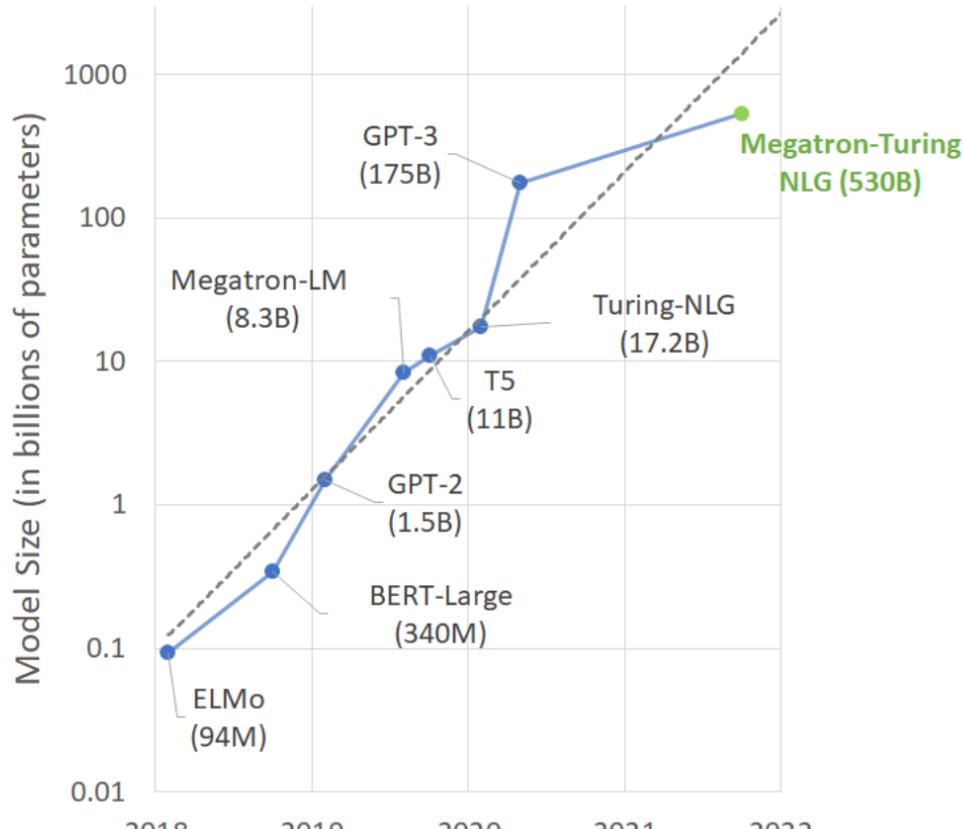


Figure 1. Trend of sizes of state-of-the-art NLP models with time

- The trend in training ever large models continues
- Gopher has **270B** parameters, trained on 300B tokens (Rae et al, 2021)
- Megatron Turing NLG has **530B** parameters, trained on 270B tokens (Smith et al., 2022)
- Chinchilla has **70B** parameters trained on **1.4T tokens (4X)** outperformed Gopher (Hoffmann et al., 2022)
- PaLM has **540B** parameters trained on 780B tokens (Chowdhery et al., 2022)

Hugging Face Open-Sourced Models

Multimodal

- Feature Extraction
- Text-to-Image
- Image-to-Text
- Text-to-Video
- Visual Question Answering
- Document Question Answering
- Graph Machine Learning

Computer Vision

- Depth Estimation
- Image Classification
- Object Detection
- Image Segmentation
- Image-to-Image
- Unconditional Image Generation
- Video Classification
- Zero-Shot Image Classification

Natural Language Processing

- Text Classification
- Token Classification
- Table Question Answering
- Question Answering
- Zero-Shot Classification
- Translation
- Summarization
- Conversational
- Text Generation
- Text2Text Generation
- Fill-Mask
- Sentence Similarity

Audio

Models 25,151

Filter

new Full-text search

↑↓ Sort: Most Download

gpt2

Text Generation • Updated Jun 29 • ↓ 22.5M • ❤ 1.4k

NousResearch/Llama-2-13b-hf

Text Generation • Updated Aug 26 • ↓ 17.9M • ❤ 43

s. stabilityai/StableBeluga-7B

Text Generation • Updated Aug 29 • ↓ 2.33M • ❤ 108

distilgpt2

Text Generation • Updated Apr 29 • ↓ 1.7M • ❤ 257

bigscience/bloom-560m

Text Generation • Updated 6 days ago • ↓ 1.35M • ❤ 255

facebook/opt-125m

Text Generation • Updated 17 days ago • ↓ 1.15M • ❤ 77

meta-llama/Llama-2-7b-chat-hf

Text Generation • Updated Aug 9 • ↓ 1.05M • ❤ 1.3k

bigscience/bloomz-1b1

Text Generation • Updated May 27 • ↓ 946k • ❤ 29

Riiid/sheep-duck-llama-2

Text Generation • Updated 26 days ago • ↓ 740k • ❤ 25

facebook/opt-1.3b

Text Generation • Updated 17 days ago • ↓ 718k • ❤ 115

TinyStories: How Small Can Language Models Be and Still Speak Coherent English?

40

- Use GPT4 generated children's stories to train small networks, Eldan and Li, 2023

Prompt	2.5M, 8 layers	33M, 4 layers	21M, 1 layer	GPT2-XL (1.5B)
Alice was so tired when she got back home so she went	to bed.	straight to bed.	to bed with a big smile on her face.	outside
Lily likes cats and dogs. She asked her mom for a dog and her mom said no, so instead she asked	her mom if she could have a dog.	her dad for a cat.	her mom again.	her dad. They brought her a Chihuahua, and she took to them immediately.
Alice and Jack walked up the street and met a girl in a red dress. The girl said to them, "Hi, I'm Jane. What are your names?"	Alice smiled and said, "My name is Daisy. What's your name?"	Alice said, "I'm Alice and this is Jack."	Jack smiled and said, "I'm Jane. Nice to meet you!"	Jane said, "Jack and Kate"